NYU

NYU Summer Machine Learning Program

Presenter Name Here
Date Here

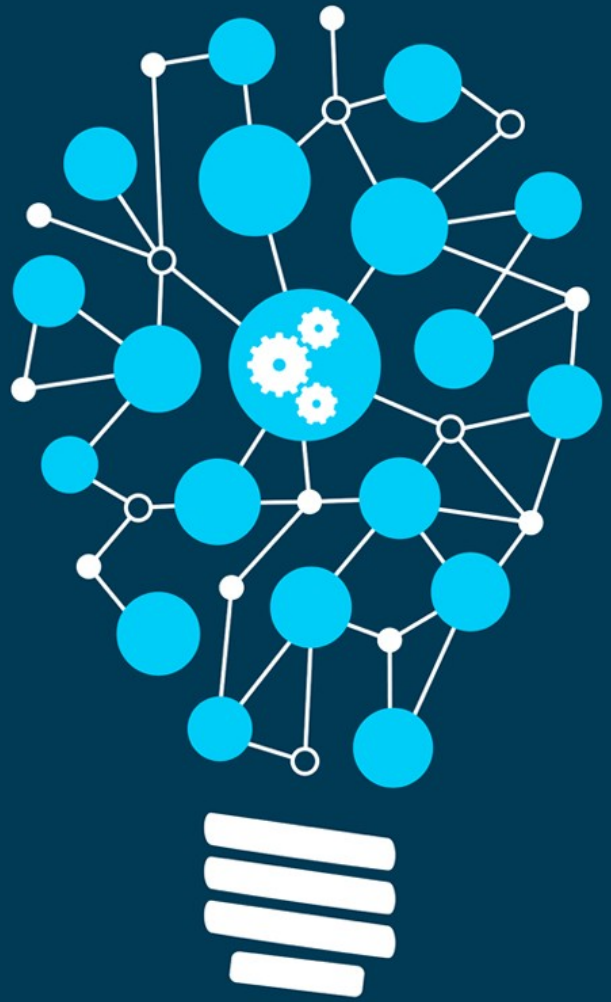

NYU

Introduction to Machine Learning

Day 1

Learning Objectives

- Fundamentals of Python programming language
- Familiarity with NumPy and Pandas
- Provide examples of Machine Learning used today
- Given a new problem, qualitatively describe how a machine learning can be used
  - Formulate a potential machine learning task
  - Identify the data needed for the task
- Classify a machine learning task
  - Regression vs. Classification
- Identify the predictors and target variables
- Determine the role of expert knowledge in the task vs. data driven learning

Outline

- Basics of Programming
  - Python loops and data structures
  - Scientific computational package – NumPy
  - Data visualization
- What is machine learning?
- Types of machine learning algorithms
  - Classification and regression
- Why the hype today?

Programming basics in Google Colab Notebook

- Google Colab is a free cloud service
  - Machine Learning education and research tool
  - Free and requires no setup
  - Supports free GPU to perform fast computations
  - You can improve your python programming skills
- Python
  - Loops
  - Data Structures
- Data Visualization
  - Load data using Pandas
  - Visualize the data by plotting histograms, scatter plots, etc.

```
[ ]  a = [20,43,6,90,78,3]
     print("forwards")
     for i in range(len(a)):
       print(a[i])

     print("now backwards:")
     i = len(a)-1
     while(i>=0):
       print(a[i])
       i -= 1
```

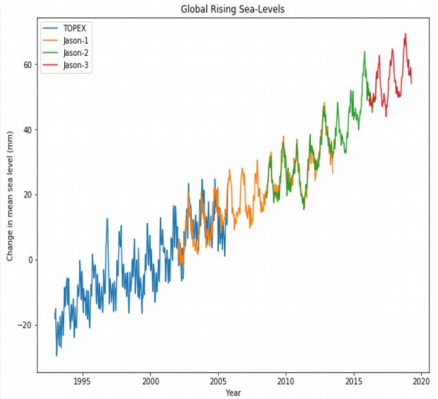


```
[ ]  a = ["apples",5,32,"oranges",10] # an example list
     a[0] # index a single element
     a[1:3] # index a slice of list (last element not included!)
     len(a) # length of the list
     b = ["bananas", a, "42"] # a list within a list
     a.append("anything"); print(a) # add element to end of list, then print the list
     # you can also remove elements. Google to find out how!
```

```
['apples', 5, 32, 'oranges', 10, 'anything']
```

Outline

- Basics of Programming
  - Python loops and data structures
  - Scientific computational package – NumPy
  - Data visualization
- What is machine learning?
- Types of machine learning algorithms
  - Classification and regression
- Why the hype today?

What is Machine Learning?

- Learn the algorithm from known data to generate the rules
- Make predictions on unknown data using these rules.

Why Machine Learning over Expert Approach?

- Human expertise does not exist (ex: complex medical processes we don't fully understand)
- Humans are unable to explain their expertise (speech recognition)
- Solution changes in time (routing on a computer network)
- Solution needs to be adapted to specific cases (user biometrics)

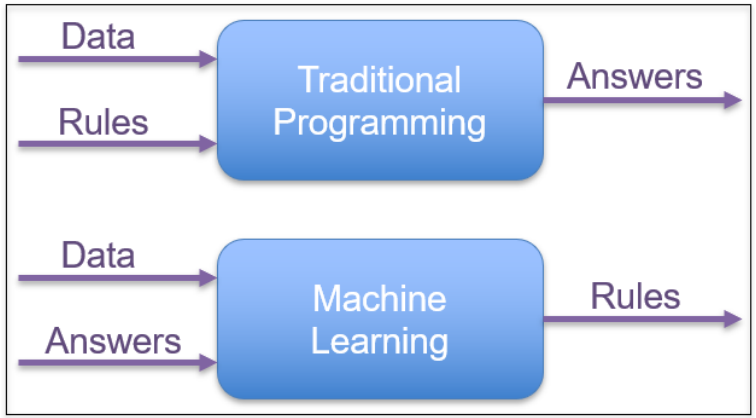

Example 1: Digit Recognition

- Problem: Recognize a digit from the image
- MNIST dataset challenge
  - Dataset developed in 1990s to spur AI research on a challenging problem for the time
  - Data taken from census forms
  - Became a classic benchmark for machine vision problems
  - We will see this dataset extensively in this class

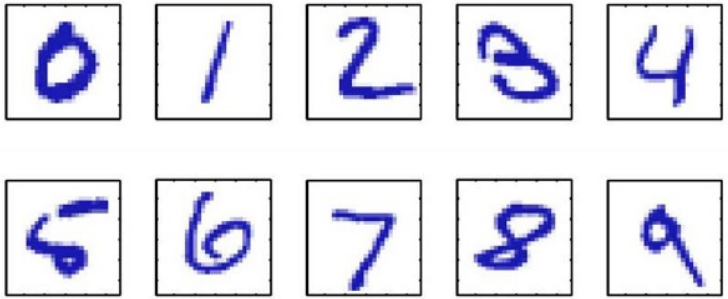

Images are 28 x 28 pixels

Example 1: Digit Recognition – Classical "Expert" Approach

- Idea: Use your knowledge about digits
  - You are an "expert" since you can do the task
  - So, you construct simple rules and code them
- Expert rule example: "Image is a digit 7 if…":
  - There is a single horizontal line, and
  - There is a single vertical line
- Rule seems simple and reasonable
- But,…

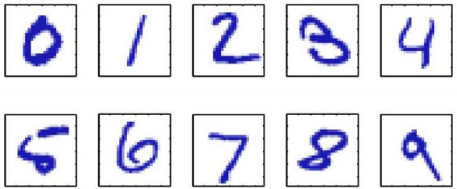

Images are 28 x 28 pixels

```
def count_vert_lines(image):
    ...
def count_horiz_lines(image):
    ...

def classify(image):
    ...
    nv = count_vert_lines(image)
    nh = count_horiz_lines(image)
    ...

    if (nv == 1) and (nh == 1):
        digit = 7
    ...

    return digit
```

Example 1: Digit Recognition – Problems with Expert Rules

- Simple expert rule breaks down in practice
  - Hard to define a “line” precisely
  - Orientation, length, thickness, …
  - May be multiple lines…
- General problem: Difficult to code our knowledge
  - We can do the task
  - But it is hard to translate to simple mathematical formula

```
def count_vert_lines(image):
    ...
def count_horiz_lines(image):
    ...

def classify(image):
    ...
    nv = count_vert_lines(image)
    nh = count_horiz_lines(image)
    ...

    if (nv == 1) and (nh == 1):
        digit = 7
    ...

    return digit
```

Example 1: Digit Recognition – Problems with Expert Rules

- Do not use your "expert" knowledge
- Learn the function from data!
- Supervised learning:
  - Get many labeled examples $(\boldsymbol{x}_i, y_i)$, i=1,…,N (Called the training data)
  - Each example has an input $\boldsymbol{x}_i$ and output y_i
  - Learn a function $f(\boldsymbol{x})$ such that: $f(\boldsymbol{x}) = y_i$ for "most" training examples

Training inputs images $\boldsymbol{x}_i$ (ex. 5000 ex per class)

?

Learned classifier
$f(\boldsymbol{x})$

Training output labels $y_i \in \{0,1,\ldots,9\}$

Example 1: Digit Recognition – ML Approach Benefits & Challenges

- Learned systems do very well on image recognition problems
  - On MNIST, current systems get <0.21% errors (as of 1/20/2018)
  - Used widely in commercial systems today (e.g. OCR)
  - Cannot match this performance with an expert system
- But there are challenges:
  - How do we acquire data? Someone must manually label examples.
  - How do we train an algorithm to learn from the data?
  - If a function works on training example, will it generalize on new data?
- This is what you will learn in this course

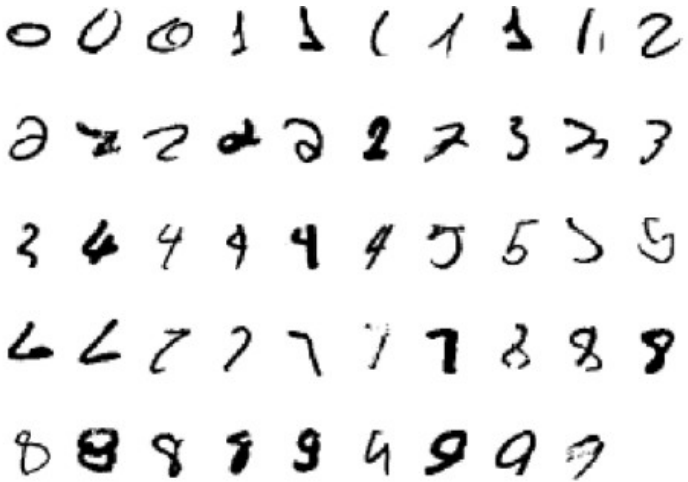

Example 2: Face Detection

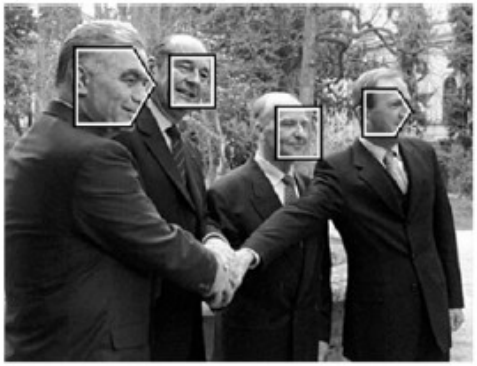

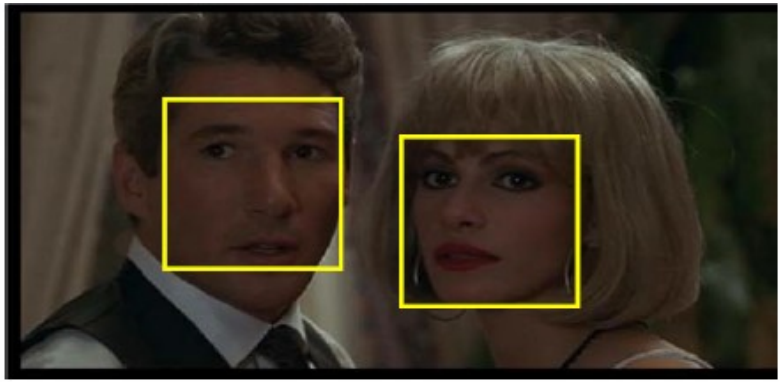

- Problem: For each image region, determine if face or non-face
- More challenging than digit recognition
  - Even harder to describe a face via “rules” in a robust way

Example 2: Face Detection - Supervised Learning Approach

- Data: Get large number of face and non-face examples
- Typical early dataset
  - 5000 faces (all near frontal, vary age, race, gender, lighting
  - 10^8 non faces
- Train an algorithm to learn the classification rules/function
  - The function maps image to binary value “face” or “non-face”
  - For good performance, functions may be complex
  - Many parameters

Example 3: Spam Detection

- Classification problem:
  - Is email junk or not junk?
- For ML, must represent email numerically
  - Common model: bag of words
  - Enumerate all words, i=1,…,N
  - Represent email via word count

    x_i = num instances of word i
- Challenge:
  - Very high-dimensional vector

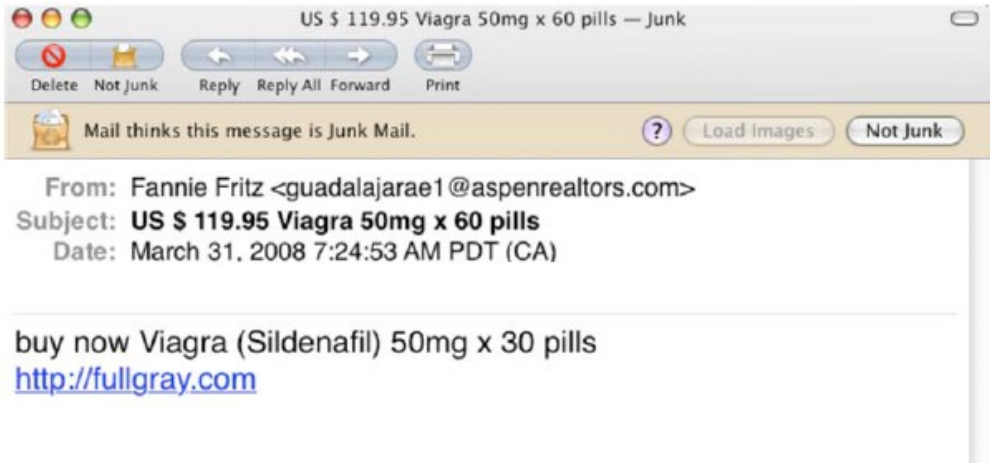

Machine Learning in Many Fields

- Retail: Market basket analysis, Customer relationship management (CRM)
- Finance: Credit scoring, fraud detection
- Manufacturing: Control, robotics, troubleshooting
- Medicine: Medical diagnosis
- Telecommunications: Spam filters, intrusion detection
- Bioinformatics: Motifs, alignment
- Web mining: Search engines

Outline

- Basics of Programming
  - Python loops and data structures
  - Scientific computational package – NumPy
  - Data visualization
- What is machine learning?
- Types of machine learning algorithms
  - Classification and regression
- Why the hype today?

Classification

- Example: Credit score
- Determine if customer is high-risk or low-risk
- Select features:
  - Example: Income & Savings
  - Represent as a vector $x=(x\_1,x\_2)$
- Learn a function from features to target
  - Use past training data
- Need to get this data

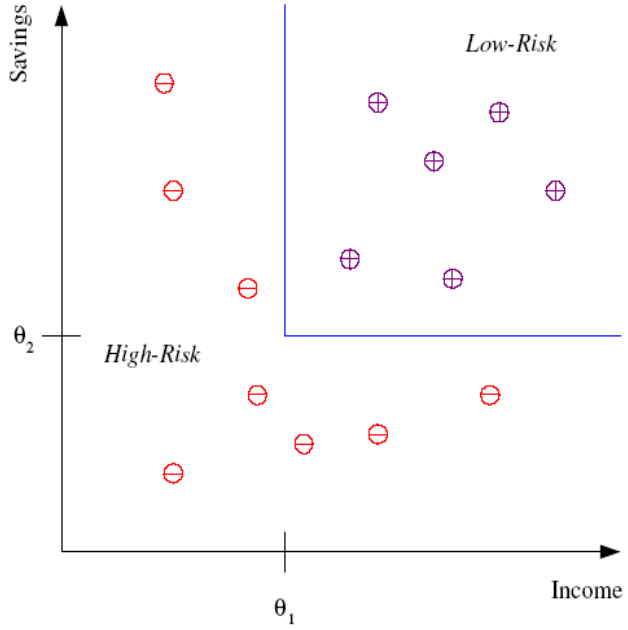

Regression

- Target variable y is continuous-valued
- Example:
  - Predict y = price of car
  - From x = mileage, size, horsepower, ..
  - Can use multiple predictors
- Assume some form of the mapping
  - Ex. Linear: $y = \beta_0 + \beta_1 * x$
  - Find parameters β_0, β_1 from data

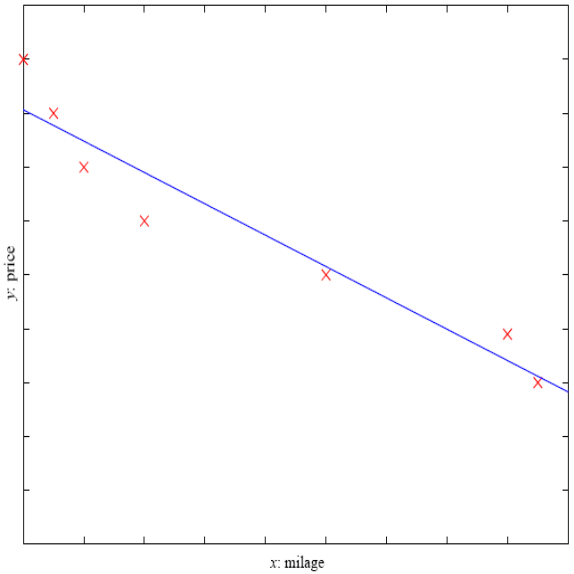

Regression Example – In Google Colab Notebook

- The Boston housing data set was collected in the 1970s
- Predict housing prices
- Many possible predictors:
  - Crime
  - Areas of non-retail business in the town
  - Age of people who own the house

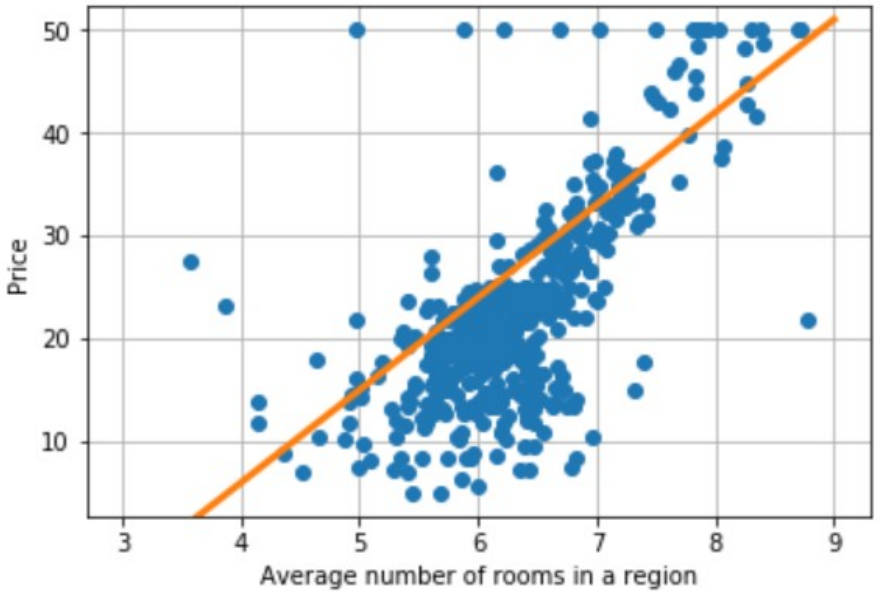

Outline

- Basics of Programming
  - Python loops and data structures
  - Scientific computational package – NumPy
  - Data visualization
- What is machine learning?
- Types of machine learning algorithms
  - Classification and regression
- Why the hype today?

What ML is Doing Today?

- Autonomous driving
- Jeopardy
- Very difficult games: Alpha Go
- Machine translation

- Many, many others…

Why Now?

- Machine learning is an old field
  - Much of the pioneering statistical work dates to the 1950s
- So what is new now?
- Big Data:
  - Massive storage. Large data centers
  - Massive connectivity
  - Sources of data from Internet and elsewhere
- Computational advances
  - Distributed machines, clusters
  - GPUs and hardware

Summary

- Fundamentals of Python programming language
- Familiarity with NumPy and Pandas
- Provide examples of Machine Learning used today
- Given a new problem, qualitatively describe how a machine learning can be used
  - Formulate a potential machine learning task
  - Identify the data needed for the task
- Classify a machine learning task
  - Regression vs. Classification
- Identify the predictors and target variables
- Determine the role of expert knowledge in the task vs. data driven learning

Thank You!