

Day 2: Linear Regression

Summer STEM: Machine Learning

Nikola Janjušević, Akshaj Kumar Veldanda, Jacky Yuan,
Tejaishwarya Gagadam

Department of Electrical Engineering
NYU Tandon School of Engineering
Brooklyn, New York

June 18, 2019

Learning Objectives

- What is the simple linear model for regression?
- What is an error function?
- What is the least squares error function?
- How do we use correlation to tell us about goodness of fit?
- What do we mean by goodness of fit?
- When is it useful to transform the target variable?
- How do we formulate the least squares problem using matrices?
- How does this extend with multi-variable features?
- What is the transformed linear model?
- What do we mean when we talk about “linear”?




Outline

- 1 Review of Day 1
- 2 Lab: Simple Linear Model
- 3 Lab: Goodness of Fit
- 4 Statistics for the LS Solution
- 5 Least Squares Solution
- 6 Extension to Multivariable Data
- 7 Lab: Robot Arm Calibration
- 8 Transforming the Output Data
- 9 Polynomial Regression
- 10 Transformed Linear Model
- 11 Lab: Fitting a Curve


What is Machine Learning

- Learn the algorithm from known data to generate the rules
- Make predictions on unknown data using these rules
- Very effective tool where human expertise is not available



Regression

- Target variable is continuous-valued
- Example
 - Predict y = price of a car
 - From x = mileage, size, horsepower
 - Can use multiple predictors
- Assume some form of mapping
 - Ex: Linear mapping:
$$y = \beta_0 + \beta_1 x$$
 - Find parameter β_0, β_1 from data
- Use target-feature pairings as examples to form model



What is Classification?

- Determine what class a target belongs to based on its features
- Example:
 - Predict y = what type of object is in a photo
 - From x = the pixels of the image
- Learn a model/function from features to target
- Use target-feature pairings as examples to form model




Outline

- 1 Review of Day 1
- 2 Lab: Simple Linear Model**
- 3 Lab: Goodness of Fit
- 4 Statistics for the LS Solution
- 5 Least Squares Solution
- 6 Extension to Multivariable Data
- 7 Lab: Robot Arm Calibration
- 8 Transforming the Output Data
- 9 Polynomial Regression
- 10 Transformed Linear Model
- 11 Lab: Fitting a Curve

Linear Model


■ Data Representation:

- y = variable you are trying to predict. Also referred to as: Dependent variable, response variable, target variable etc.
- x = what you are using to predict. Also referred to as: Independent variable, attribute, predictor etc.
- Set of points, (x_i, y_i) , $i = 1, \dots, n$. Each data point is called a sample.
- An efficient way to visualize the data is by plotting y vs x in a scatter plot.



Linear Model

- Assume a linear relationship $y = \beta_0 + \beta_1 x$
 - β_0 = intercept
 - β_1 = slope
- $\beta = (\beta_0, \beta_1)$ are the parameters of the model



- Let's go to the lab to understand this further.

Outline

- 1 Review of Day 1
- 2 Lab: Simple Linear Model
- 3 Lab: Goodness of Fit**
- 4 Statistics for the LS Solution
- 5 Least Squares Solution
- 6 Extension to Multivariable Data
- 7 Lab: Robot Arm Calibration
- 8 Transforming the Output Data
- 9 Polynomial Regression
- 10 Transformed Linear Model
- 11 Lab: Fitting a Curve

Is Your Model a Good Fit?

- How would you determine if your model is a good fit or not?
- Talk with your classmates next to you to see whose model fits the data the best
 - How will you determine this?
 - Is there a quantitative way?
 - Write python code if so.

Error Functions

- An **error function** quantifies the discrepancy between your model and the data.
 - They are non-negative, and go to zero as the model gets better.
- Common Error Functions:
 - Mean Squared Error: $MSE = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y})^2$
 - Mean Absolute Error: $MAE = \frac{1}{N} \sum_{i=1}^N |y_i - \hat{y}|$
- In later units, we will refer to these as **cost functions** or **loss functions**.
- Compute MSE on your model

General Steps to Solve a Machine Learning Problem

- Load and visualize data

General Steps to Solve a Machine Learning Problem

- Load and visualize data
 - $(x_i, y_i), i = 1, \dots, n$

General Steps to Solve a Machine Learning Problem

- Load and visualize data
 - $(x_i, y_i), i = 1, \dots, n$
- Find an appropriate model to fit the data

General Steps to Solve a Machine Learning Problem

- Load and visualize data
 - $(x_i, y_i), i = 1, \dots, n$
- Find an appropriate model to fit the data
 - Eg: Linear model is $\hat{y} = \beta_1 x + \beta_0$

General Steps to Solve a Machine Learning Problem

- Load and visualize data
 - $(x_i, y_i), i = 1, \dots, n$
- Find an appropriate model to fit the data
 - Eg: Linear model is $\hat{y} = \beta_1 x + \beta_0$
- Choose an appropriate error function

General Steps to Solve a Machine Learning Problem

- Load and visualize data
 - $(x_i, y_i), i = 1, \dots, n$
- Find an appropriate model to fit the data
 - Eg: Linear model is $\hat{y} = \beta_1 x + \beta_0$
- Choose an appropriate error function
 - $MSE = \sum_{i=1}^N (y_i - (\beta_0 + \beta_1 x_i))^2$

General Steps to Solve a Machine Learning Problem

- Load and visualize data
 - $(x_i, y_i), i = 1, \dots, n$
- Find an appropriate model to fit the data
 - Eg: Linear model is $\hat{y} = \beta_1 x + \beta_0$
- Choose an appropriate error function
 - $MSE = \sum_{i=1}^N (y_i - (\beta_0 + \beta_1 x_i))^2$
- Find parameters that minimize the error function

General Steps to Solve a Machine Learning Problem

- Load and visualize data
 - $(x_i, y_i), i = 1, \dots, n$
- Find an appropriate model to fit the data
 - Eg: Linear model is $\hat{y} = \beta_1 x + \beta_0$
- Choose an appropriate error function
 - $MSE = \sum_{i=1}^N (y_i - (\beta_0 + \beta_1 x_i))^2$
- Find parameters that minimize the error function
 - Select β_0, β_1 to minimize the error function

Least Squares Fit

- The **Least Squares Fit** is characterized by the minimization of the MSE error function:

$$MSE = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y})^2$$

Least Squares Fit

- The **Least Squares Fit** is characterized by the minimization of the MSE error function:

$$MSE = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y})^2$$

- Find the parameters, $\beta = (\beta_0, \beta_1)$, that give the smallest MSE

Least Squares Fit

- The **Least Squares Fit** is characterized by the minimization of the MSE error function:

$$MSE = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y})^2$$

- Find the parameters, $\beta = (\beta_0, \beta_1)$, that give the smallest MSE
- MSE is a useful metric because there exists an analytic solution to find the optimal parameters β_0 and β_1

Outline

- 1 Review of Day 1
- 2 Lab: Simple Linear Model
- 3 Lab: Goodness of Fit
- 4 Statistics for the LS Solution**
- 5 Least Squares Solution
- 6 Extension to Multivariable Data
- 7 Lab: Robot Arm Calibration
- 8 Transforming the Output Data
- 9 Polynomial Regression
- 10 Transformed Linear Model
- 11 Lab: Fitting a Curve

Mean, Variance, and Covariance, Correlation Coefficient

- Given feature-target data
 $(x_i, y_i), i = 1, 2, \dots, N$

Mean, Variance, and Covariance, Correlation Coefficient

- Given feature-target data
 $(x_i, y_i), i = 1, 2, \dots, N$
- Mean:

$$\bar{x} = \frac{1}{N} \sum_{i=1}^N x_i, \quad \bar{y} = \frac{1}{N} \sum_{i=1}^N y_i$$

Mean, Variance, and Covariance, Correlation Coefficient

- Given feature-target data
 $(x_i, y_i), i = 1, 2, \dots, N$

- Mean:

$$\bar{x} = \frac{1}{N} \sum_{i=1}^N x_i, \quad \bar{y} = \frac{1}{N} \sum_{i=1}^N y_i$$

- Variance:

$$\sigma_x^2 = \frac{1}{N} \sum_{i=1}^N (x_i - \bar{x})^2, \quad \sigma_y^2 = \frac{1}{N} \sum_{i=1}^N (y_i - \bar{y})^2$$

Mean, Variance, and Covariance, Correlation Coefficient

- Given feature-target data (x_i, y_i) , $i = 1, 2, \dots, N$

- Mean:

$$\bar{x} = \frac{1}{N} \sum_{i=1}^N x_i, \quad \bar{y} = \frac{1}{N} \sum_{i=1}^N y_i$$

- Variance:

$$\sigma_x^2 = \frac{1}{N} \sum_{i=1}^N (x_i - \bar{x})^2, \quad \sigma_y^2 = \frac{1}{N} \sum_{i=1}^N (y_i - \bar{y})^2$$

- **Covariance:**

$$\sigma_{xy} = \frac{1}{N} \sum_{i=1}^N (x_i - \bar{x})(y_i - \bar{y})$$

Mean, Variance, and Covariance, Correlation Coefficient

- Given feature-target data (x_i, y_i) , $i = 1, 2, \dots, N$

- Mean:

$$\bar{x} = \frac{1}{N} \sum_{i=1}^N x_i, \quad \bar{y} = \frac{1}{N} \sum_{i=1}^N y_i$$

- Variance:

$$\sigma_x^2 = \frac{1}{N} \sum_{i=1}^N (x_i - \bar{x})^2, \quad \sigma_y^2 = \frac{1}{N} \sum_{i=1}^N (y_i - \bar{y})^2$$


- **Covariance:**

$$\sigma_{xy} = \frac{1}{N} \sum_{i=1}^N (x_i - \bar{x})(y_i - \bar{y})$$

- **Correlation Coefficient:**

$$\rho = \frac{\sigma_{xy}}{\sigma_x \sigma_y}$$

Lab: Gaining Intuition



Outline

- 1 Review of Day 1
- 2 Lab: Simple Linear Model
- 3 Lab: Goodness of Fit
- 4 Statistics for the LS Solution
- 5 Least Squares Solution**
- 6 Extension to Multivariable Data
- 7 Lab: Robot Arm Calibration
- 8 Transforming the Output Data
- 9 Polynomial Regression
- 10 Transformed Linear Model
- 11 Lab: Fitting a Curve

LS Fit Solution

■ Model:

$$\hat{y} = \beta_0 + \beta_1 x$$

LS Fit Solution

- Model:

$$\hat{y} = \beta_0 + \beta_1 x$$

- Optimization:

$$MSE = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y})^2$$



LS Fit Solution

- Model:

$$\hat{y} = \beta_0 + \beta_1 x$$

- Optimization:

$$MSE = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y})^2$$

- Solution:

$$\hat{y} = \bar{y} + \rho \frac{\sigma_y}{\sigma_x} (x - \bar{x})$$

LS Fit Solution

- Model:

$$\hat{y} = \beta_0 + \beta_1 x$$

- Optimization:

$$MSE = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y})^2$$

- Solution:

$$\hat{y} = \bar{y} + \rho \frac{\sigma_y}{\sigma_x} (x - \bar{x})$$

$$\beta_1 = \rho \frac{\sigma_y}{\sigma_x}, \quad \beta_0 = \bar{y} - \beta_1 \bar{x}$$

LS Fit Solution

- Model:

$$\hat{y} = \beta_0 + \beta_1 x$$

- Optimization:

$$MSE = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y})^2$$

- Solution:

$$\hat{y} = \bar{y} + \rho \frac{\sigma_y}{\sigma_x} (x - \bar{x})$$

$$\beta_1 = \rho \frac{\sigma_y}{\sigma_x}, \quad \beta_0 = \bar{y} - \beta_1 \bar{x}$$

- Prediction:

$$y_{new} = \beta_0 + \beta_1 x_{new}$$

LS Fit Solution

- Model:

$$\hat{y} = \beta_0 + \beta_1 x$$

- Optimization:

$$MSE = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y})^2$$

- Solution:

$$\hat{y} = \bar{y} + \rho \frac{\sigma_y}{\sigma_x} (x - \bar{x})$$

$$\beta_1 = \rho \frac{\sigma_y}{\sigma_x}, \quad \beta_0 = \bar{y} - \beta_1 \bar{x}$$

- Prediction:

$$y_{new} = \beta_0 + \beta_1 x_{new}$$

- Compute the LS fit model

Lab: Find/Build and fit your own data

1 Find your a data set

- Google: “[subject you’re interested in] dataset”
- <https://archive.ics.uci.edu/ml/datasets.php>
- <https://toolbox.google.com/datasetsearch>

– or –

2 Build your own data set

- Examples:
 - Use Number of CPUs in a Computer to Predict the Price
 - Number of Instagram Followers to predict likes
- Only need 10-15 samples

Outline

- 1 Review of Day 1
- 2 Lab: Simple Linear Model
- 3 Lab: Goodness of Fit
- 4 Statistics for the LS Solution
- 5 Least Squares Solution
- 6 Extension to Multivariable Data**
- 7 Lab: Robot Arm Calibration
- 8 Transforming the Output Data
- 9 Polynomial Regression
- 10 Transformed Linear Model
- 11 Lab: Fitting a Curve

Extending the Model

■ Model: $\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_D x_D$

Extending the Model

- Model: $\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_D x_D$
- Suppose we have multiple features corresponding to a single target output:

Extending the Model

- Model: $\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_D x_D$
- Suppose we have multiple features corresponding to a single target output:
 - y , single target variable

Extending the Model

- Model: $\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_D x_D$
- Suppose we have multiple features corresponding to a single target output:
 - y , single target variable
 - multiple features: x_1, x_2, \dots, x_D

Extending the Model

- Model: $\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_D x_D$
- Suppose we have multiple features corresponding to a single target output:
 - y , single target variable
 - multiple features: x_1, x_2, \dots, x_D
 - Target variable is a linear combination of all feature variables



Extending the Model

- Model: $\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_D x_D$
- Suppose we have multiple features corresponding to a single target output:
 - y , single target variable
 - multiple features: x_1, x_2, \dots, x_D
 - Target variable is a linear combination of all feature variables
 - **Linear Combination**: sum of scaled features



Extending the Model

- Model: $\hat{y} = \beta_0 \times 1 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_D x_D$
- Suppose we have multiple features corresponding to a single target output:
 - y , single target variable
 - multiple features: x_1, x_2, \dots, x_D
 - Target variable is a linear combination of all feature variables
 - **Linear Combination**: sum of scaled features
- Vector Notation:

Extending the Model

- Model: $\hat{y} = \beta_0 \times 1 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_D x_D$
- Suppose we have multiple features corresponding to a single target output:
 - y , single target variable
 - multiple features: x_1, x_2, \dots, x_D
 - Target variable is a linear combination of all feature variables
 - **Linear Combination**: sum of scaled features
- Vector Notation:
Let $\mathbf{x} = [1, x_1, x_2, \dots, x_D]^T$

Extending the Model

- Model: $\hat{y} = \beta_0 \times 1 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_D x_D$
- Suppose we have multiple features corresponding to a single target output:
 - y , single target variable
 - multiple features: x_1, x_2, \dots, x_D
 - Target variable is a linear combination of all feature variables
 - **Linear Combination**: sum of scaled features
- Vector Notation:
 - Let $\mathbf{x} = [1, x_1, x_2, \dots, x_D]^T$
 - and $\boldsymbol{\beta} = [\beta_0, \beta_1, \dots, \beta_D]^T$

Extending the Model

- Model: $\hat{y} = \beta_0 \times 1 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_D x_D$
- Suppose we have multiple features corresponding to a single target output:
 - y , single target variable
 - multiple features: x_1, x_2, \dots, x_D
 - Target variable is a linear combination of all feature variables
 - **Linear Combination**: sum of scaled features
- Vector Notation:
 - Let $\mathbf{x} = [1, x_1, x_2, \dots, x_D]^T$
 - and $\boldsymbol{\beta} = [\beta_0, \beta_1, \dots, \beta_D]^T$
 - So, $\hat{y} = \mathbf{x}^T \boldsymbol{\beta}$

Matrix Formulation

- We want to minimize the squared error over all the samples:

$$\sum_{i=1}^N (y_i - \hat{y}_i)^2 = \sum_{i=1}^N (y_i - \mathbf{x}_i^T \boldsymbol{\beta})^2$$

Matrix Formulation

- We want to minimize the squared error over all the samples:

$$\sum_{i=1}^N (y_i - \hat{y}_i)^2 = \sum_{i=1}^N (y_i - \mathbf{x}_i^T \boldsymbol{\beta})^2$$

- Design Matrix: Let, $A = \begin{bmatrix} 1 & x_{1_1} & \cdots & x_{1_D} \\ 1 & x_{2_1} & \cdots & x_{2_D} \\ \vdots & & \ddots & \\ 1 & x_{N_1} & \cdots & x_{N_D} \end{bmatrix}$



Matrix Formulation

- We want to minimize the squared error over all the samples:

$$\sum_{i=1}^N (y_i - \hat{y}_i)^2 = \sum_{i=1}^N (y_i - \mathbf{x}_i^T \boldsymbol{\beta})^2$$

- Design Matrix: Let, $A = \begin{bmatrix} 1 & x_{1_1} & \cdots & x_{1_D} \\ 1 & x_{2_1} & \cdots & x_{2_D} \\ \vdots & & \ddots & \\ 1 & x_{N_1} & \cdots & x_{N_D} \end{bmatrix}$

- Can express these conditions in matrix form: $\mathbf{y} = A\boldsymbol{\beta}$



Matrix Formulation

- We want to minimize the squared error over all the samples:

$$\sum_{i=1}^N (y_i - \hat{y}_i)^2 = \sum_{i=1}^N (y_i - \mathbf{x}_i^T \boldsymbol{\beta})^2$$

- Design Matrix: Let, $A = \begin{bmatrix} 1 & x_{1_1} & \cdots & x_{1_D} \\ 1 & x_{2_1} & \cdots & x_{2_D} \\ \vdots & & \ddots & \\ 1 & x_{N_1} & \cdots & x_{N_D} \end{bmatrix}$

- Can express these conditions in matrix form: $\mathbf{y} = A\boldsymbol{\beta}$
- Solution: Pseudo-Inverse

Matrix Formulation

- We want to minimize the squared error over all the samples:

$$\sum_{i=1}^N (y_i - \hat{y}_i)^2 = \sum_{i=1}^N (y_i - \mathbf{x}_i^T \boldsymbol{\beta})^2$$

- Design Matrix: Let, $A = \begin{bmatrix} 1 & x_{1_1} & \cdots & x_{1_D} \\ 1 & x_{2_1} & \cdots & x_{2_D} \\ \vdots & & \ddots & \\ 1 & x_{N_1} & \cdots & x_{N_D} \end{bmatrix}$

- Can express these conditions in matrix form: $\mathbf{y} = A\boldsymbol{\beta}$
- Solution: Pseudo-Inverse
 - There isn't a true solution to this equation.

Matrix Formulation

- We want to minimize the squared error over all the samples:

$$\sum_{i=1}^N (y_i - \hat{y}_i)^2 = \sum_{i=1}^N (y_i - \mathbf{x}_i^T \boldsymbol{\beta})^2$$

- Design Matrix: Let, $A = \begin{bmatrix} 1 & x_{1_1} & \cdots & x_{1_D} \\ 1 & x_{2_1} & \cdots & x_{2_D} \\ \vdots & & \ddots & \\ 1 & x_{N_1} & \cdots & x_{N_D} \end{bmatrix}$

- Can express these conditions in matrix form: $\mathbf{y} = A\boldsymbol{\beta}$

- Solution: Pseudo-Inverse

- There isn't a true solution to this equation.

- We say $\boldsymbol{\beta}^*$ solves $\mathbf{y} = A\boldsymbol{\beta}$ in the least squares sense, where

$$\boldsymbol{\beta}^* = A^\dagger \mathbf{y}$$

Boston Housing Demo

- Using multiple features to predicting house prices
 - Crime rate per capita, number of rooms, student-teacher ratio at local schools, ...
- Lets use a linear model that takes into account all the collected data
- Go to the Github, Day2/Day2.ipynb

M6: Demo: Multivariable Regression on Boston Housing Data

```
[ ] import pandas as pd
import numpy as np

df = pd.read_csv('https://archive.ics.uci.edu/ml/machine-learning-databases/housing/housing.data',
                 header=None, delin_whitespace=True, names=names, na_values='?')

df.head(6) # print the first six samples
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTSTRATIO	B	LSTAT	PRICE
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296.0	15.3	396.90	4.98	24.0
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242.0	17.8	396.90	9.14	21.6
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242.0	17.8	392.83	4.03	34.7
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222.0	18.7	394.63	2.94	33.4
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222.0	18.7	396.90	5.33	36.2
5	0.02985	0.0	2.18	0	0.458	6.430	58.7	6.0622	3	222.0	18.7	394.12	5.21	28.7

Outline

- 1 Review of Day 1
- 2 Lab: Simple Linear Model
- 3 Lab: Goodness of Fit
- 4 Statistics for the LS Solution
- 5 Least Squares Solution
- 6 Extension to Multivariable Data
- 7 Lab: Robot Arm Calibration**
- 8 Transforming the Output Data
- 9 Polynomial Regression
- 10 Transformed Linear Model
- 11 Lab: Fitting a Curve

Robot Arm Calibration


- Let's train a model based on the given data.
- In this lab we're going to:
 - Predict the *current* drawn
 - Predictors, X : Robot arm's joint angles, velocity, acceleration, strain gauge readings (load measurement).

Outline

- 1 Review of Day 1
- 2 Lab: Simple Linear Model
- 3 Lab: Goodness of Fit
- 4 Statistics for the LS Solution
- 5 Least Squares Solution
- 6 Extension to Multivariable Data
- 7 Lab: Robot Arm Calibration
- 8 Transforming the Output Data**
- 9 Polynomial Regression
- 10 Transformed Linear Model
- 11 Lab: Fitting a Curve

Transforming the Output Data

- Not all data can be modeled using linear relation:
 - $\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_D x_D$
- Modeling nonlinear data with linear function does not result in a good fit
- We can transform output with a nonlinear function
- What transformation we use depends on the nature of the data
- For example: We might use an exponential model for:
 - Radioactive decay
 - Population growth



- Demo: mpg of cars data

Outline

- 1 Review of Day 1
- 2 Lab: Simple Linear Model
- 3 Lab: Goodness of Fit
- 4 Statistics for the LS Solution
- 5 Least Squares Solution
- 6 Extension to Multivariable Data
- 7 Lab: Robot Arm Calibration
- 8 Transforming the Output Data
- 9 Polynomial Regression**
- 10 Transformed Linear Model
- 11 Lab: Fitting a Curve

Polynomial Regression


- What if our data is clearly *not* linear, but instead follows a polynomial curve?
 - Ex: Projectile motion, gravity, Coulomb's law, ...

Polynomial Regression

- What if our data is clearly *not* linear, but instead follows a polynomial curve?
 - Ex: Projectile motion, gravity, Coulomb's law, ...
- Can we perform regression to get a polynomial model?


Polynomial Regression

- What if our data is clearly *not* linear, but instead follows a polynomial curve?
 - Ex: Projectile motion, gravity, Coulomb's law, ...
- Can we perform regression to get a polynomial model?
 - Could it be linear regression? (What do we mean by linear?)




Polynomial Regression

- What if our data is clearly *not* linear, but instead follows a polynomial curve?
 - Ex: Projectile motion, gravity, Coulomb's law, ...
- Can we perform regression to get a polynomial model?
 - Could it be linear regression? (What do we mean by linear?)
 - What would our features be?




Polynomial Regression

- What if our data is clearly *not* linear, but instead follows a polynomial curve?
 - Ex: Projectile motion, gravity, Coulomb's law, ...
- Can we perform regression to get a polynomial model?
 - Could it be linear regression? (What do we mean by linear?)
 - What would our features be?
- Polynomial Model: $\hat{y} = \sum_{i=0}^N \beta_i x^i$




Polynomial Regression

- What if our data is clearly *not* linear, but instead follows a polynomial curve?
 - Ex: Projectile motion, gravity, Coulomb's law, ...
- Can we perform regression to get a polynomial model?
 - Could it be linear regression? (What do we mean by linear?)
 - What would our features be?
- Polynomial Model: $\hat{y} = \sum_{i=0}^N \beta_i x^i$
 - Is this linear?



Polynomial Regression

- What if our data is clearly *not* linear, but instead follows a polynomial curve?
 - Ex: Projectile motion, gravity, Coulomb's law, ...
- Can we perform regression to get a polynomial model?
 - Could it be linear regression? (What do we mean by linear?)
 - What would our features be?
- Polynomial Model: $\hat{y} = \sum_{i=0}^N \beta_i x^i$
 - Is this linear?
- Demo on Github



Outline

- 1 Review of Day 1
- 2 Lab: Simple Linear Model
- 3 Lab: Goodness of Fit
- 4 Statistics for the LS Solution
- 5 Least Squares Solution
- 6 Extension to Multivariable Data
- 7 Lab: Robot Arm Calibration
- 8 Transforming the Output Data
- 9 Polynomial Regression
- 10 Transformed Linear Model**
- 11 Lab: Fitting a Curve

Transformed Linear Model

- We can extend polynomial fitting to a more general model
- In polynomial fitting, we used the equation:
 - $y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_D x_D$
- In the more general model, the output is a linear combination of transformed input
 - $y = \beta_0 + \beta_1 \phi_1(x) + \beta_2 \phi_2(x) + \dots + \beta_D \phi_D(x)$
 - where $\phi_1(x), \phi_2(x) \dots, \phi_D(x)$ are called **basis** functions
 - Polynomial fitting is a special case where the basis functions are power functions
- Besides polynomials, we can also use other function as our basis function
 - Gaussian: $\phi(x) = e^{-\frac{(x-\mu)^2}{2\sigma^2}}$
 - Exponential: $\phi(x) = e^{-\alpha x}$



Transformed Linear Model

- How do we fit a transformed linear model?
 - $y = \beta_0 + \beta_1\phi_1(\mathbf{x}) + \beta_2\phi_2(\mathbf{x}) + \dots + \beta_D\phi_D(\mathbf{x})$
- The procedure is similar to polynomial fitting
- First transform the features of each example using the transformation

- Form the Design Matrix: $\Phi = \begin{bmatrix} \phi_0(\mathbf{x}) & \dots & \phi_D(\mathbf{x}) \\ \vdots & \ddots & \vdots \\ \phi_0(\mathbf{x}) & \dots & \phi_D(\mathbf{x}) \end{bmatrix}$

- Solve for the Least-squares solution:
 - $\beta = \Phi^\dagger y$
 - Model: $\hat{y} = \phi(\mathbf{x})^T \beta$

Transformed Linear Model

- When the input data has multiple features, the \mathbf{x} in the previous equation,


$y = \beta_0 + \beta_1\phi_1(\mathbf{x}) + \beta_2\phi_2(\mathbf{x}) + \dots + \beta_D\phi_D(\mathbf{x})$, can also be regarded as a vector

- The transformations are then multivariate functions that uses multiple features in generating each new feature
- One example would be the multivariate Gaussian function:

- $\phi(\mathbf{x}_1, \mathbf{x}_2) = e^{-\frac{(\mathbf{x}_1 - \mu_1)^2 + (\mathbf{x}_2 - \mu_2)^2}{2\sigma^2}}$ Similar to the 1D Gaussian function which has a bell shaped curve centered at the mean, the 2D Gaussian function is a 2D bell shaped curve centered at (μ_1, μ_2)

Transformed Linear Model

■ Shape of a 2D Gaussian function



Outline

- 1 Review of Day 1
- 2 Lab: Simple Linear Model
- 3 Lab: Goodness of Fit
- 4 Statistics for the LS Solution
- 5 Least Squares Solution
- 6 Extension to Multivariable Data
- 7 Lab: Robot Arm Calibration
- 8 Transforming the Output Data
- 9 Polynomial Regression
- 10 Transformed Linear Model
- 11 Lab: Fitting a Curve**

Lab: Fitting a curve with Transformed Features

- Open Lab Notebook
- Do the lab in Module 11
- From the plot, think about what function can you use to transform the feature and perform a regression

Learning Objectives

- What is the simple linear model for regression?
- What is an error function?
- What is the least squares error function?
- How do we use correlation to tell us about goodness of fit?
- What do we mean by goodness of fit?
- When is it useful to transform the target variable?
- How do we formulate the least squares problem using matrices?
- How does this extend with multi-variable features?
- What is the transformed linear model?
- What do we mean when we talk about “linear”?

Thank You!

- Next Class: Generalization Error
- How do our models hold up against prediction new data?