

Day 3: Generalization Error

Summer STEM: Machine Learning

Nikola Janjušević, Akshaj Kumar Veldanda, Jacky Yuan,
Tejaishwarya Gagadam

Department of Electrical Engineering
NYU Tandon School of Engineering
Brooklyn, New York

June 19, 2019

Learning Objectives

- What is the difference between train error and test error?
- What is overfitting? How do we detect it?
- What is cross validation?
- How to find the optimal model order for my model?
- What is regularization? How does it prevent overfitting?

Outline

- 1 Review of Day 2
- 2 Lab: Robot Arm Calibration
- 3 Polynomial Regression
- 4 Train and Test Error, Overfitting
- 5 Model Order Selection
- 6 Regularization

General Steps to Solve a Linear Regression Problem

- Load and visualize data

General Steps to Solve a Linear Regression Problem

- Load and visualize data
 - $(x_i, y_i), i = 1, \dots, n$

General Steps to Solve a Linear Regression Problem

- Load and visualize data
 - $(x_i, y_i), i = 1, \dots, n$
- Find an appropriate model to fit the data

General Steps to Solve a Linear Regression Problem

- Load and visualize data
 - $(x_i, y_i), i = 1, \dots, n$
- Find an appropriate model to fit the data
 - Eg: Linear model is $\hat{y} = \beta_1 x + \beta_0$

General Steps to Solve a Linear Regression Problem

- Load and visualize data
 - $(x_i, y_i), i = 1, \dots, n$
- Find an appropriate model to fit the data
 - Eg: Linear model is $\hat{y} = \beta_1 x + \beta_0$
- Choose an appropriate error function

General Steps to Solve a Linear Regression Problem

- Load and visualize data
 - $(x_i, y_i), i = 1, \dots, n$
- Find an appropriate model to fit the data
 - Eg: Linear model is $\hat{y} = \beta_1 x + \beta_0$
- Choose an appropriate error function
 - $MSE = \sum_{i=1}^N (y_i - (\beta_0 + \beta_1 x_i))^2$

General Steps to Solve a Linear Regression Problem

- Load and visualize data
 - $(x_i, y_i), i = 1, \dots, n$
- Find an appropriate model to fit the data
 - Eg: Linear model is $\hat{y} = \beta_1 x + \beta_0$
- Choose an appropriate error function
 - $MSE = \sum_{i=1}^N (y_i - (\beta_0 + \beta_1 x_i))^2$
- Find parameters that minimize the error function

General Steps to Solve a Linear Regression Problem

- Load and visualize data
 - $(x_i, y_i), i = 1, \dots, n$
- Find an appropriate model to fit the data
 - Eg: Linear model is $\hat{y} = \beta_1 x + \beta_0$
- Choose an appropriate error function
 - $MSE = \sum_{i=1}^N (y_i - (\beta_0 + \beta_1 x_i))^2$
- Find parameters that minimize the error function
 - Select β_0, β_1 to minimize the error function
$$\beta_1 = \rho \frac{\sigma_y}{\sigma_x}, \quad \beta_0 = \bar{y} - \beta_1 \bar{x}$$

General Steps to Solve a Linear Regression Problem

- Load and visualize data
 - $(x_i, y_i), i = 1, \dots, n$
- Find an appropriate model to fit the data
 - Eg: Linear model is $\hat{y} = \beta_1 x + \beta_0$
- Choose an appropriate error function
 - $MSE = \sum_{i=1}^N (y_i - (\beta_0 + \beta_1 x_i))^2$
- Find parameters that minimize the error function
 - Select β_0, β_1 to minimize the error function
$$\beta_1 = \rho \frac{\sigma_y}{\sigma_x}, \quad \beta_0 = \bar{y} - \beta_1 \bar{x}$$
- Prediction: $y_{new} = \beta_0 + \beta_1 x_{new}$

Extending the Model to Multi-variable Data

- Model: $\hat{y} = \beta_0 \times 1 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_D x_D$

- Design Matrix: Let, $A = \begin{bmatrix} 1 & x_{1_1} & \cdots & x_{1_D} \\ 1 & x_{2_1} & \cdots & x_{2_D} \\ \vdots & & \ddots & \\ 1 & x_{N_1} & \cdots & x_{N_D} \end{bmatrix}$

- We say β^* solves $\mathbf{y} = A\beta$ in the least squares sense, where

$$\beta^* = A^\dagger \mathbf{y}$$

- This β^* minimizes the mean squared error

Outline

- 1 Review of Day 2
- 2 Lab: Robot Arm Calibration
- 3 Polynomial Regression
- 4 Train and Test Error, Overfitting
- 5 Model Order Selection
- 6 Regularization

Robot Arm Calibration


- Let's train a model based on the given data.
- In this lab we're going to:
 - Predict the *current* drawn
 - Predictors, X : Robot arm's joint angles, velocity, acceleration, strain gauge readings (load measurement).

Outline

- 1 Review of Day 2
- 2 Lab: Robot Arm Calibration
- 3 Polynomial Regression
- 4 Train and Test Error, Overfitting
- 5 Model Order Selection
- 6 Regularization

Polynomial Fitting


- We have been using linear model to fit our data. But it doesn't work well every time
- Some data have more complex relation that cannot be fitted well using a straight line
 - Ex: Projectile motion, Coulomb's law, Exponential growth/decay, ...



- Linear model does not look like a good fit for this data
- Can we use some other model to fit this data?

Polynomial Fitting

- Can we use a polynomial to fit our data?
- Polynomial: A sum of different powers of a variable
 - Examples: $y = x^2 + 2$, $y = 5x^3 - 3x^2 + 4$



- Polynomial Model: $y = \beta_0 + \beta_1x + \beta_2x^2 + \beta_3x^3 + \dots$

Polynomial Fitting

- Polynomial Model: $y = \beta_0 + \beta_1x + \beta_2x^2 + \beta_3x^3 + \dots$
- The process of fitting a polynomial is similar to linearly fitting multivariate data
- Recall the linear model for multivariable
- $y = \beta_0 + \beta_1x_1 + \beta_2x_2 + \beta_3x_3 + \dots$
 - Where $x_1, x_2, x_3 \dots$ are different features
- If we treat x^2 as our second feature, x^3 as our third feature, x^4 as our fourth feature.... We can use the same procedure in multivariate regression for linear fit!

Polynomial Fitting

- Design Matrix for Linear:

$$A = \begin{bmatrix} 1 & x_1 \\ 1 & x_2 \\ \vdots & \vdots \\ 1 & x_N \end{bmatrix}$$


- Design Matrix for Polynomial: $A =$

$$\begin{bmatrix} 1 & x_1 & x_1^2 & \cdots & x_1^D \\ 1 & x_2 & x_2^2 & \cdots & x_2^D \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_N & x_N^2 & \cdots & x_N^D \end{bmatrix}$$

- For the polynomial fitting, we just added columns of features that are powers of the original feature

Lab: Fit a polynomial

- You are given the data set below with x and y values



- Try to fit the data using a polynomial with a certain degree
- Calculate mean square error between the sample y and your predicted y
- Try different polynomial degree and see if you can improve the mse
- Plot your polynomial over the data points

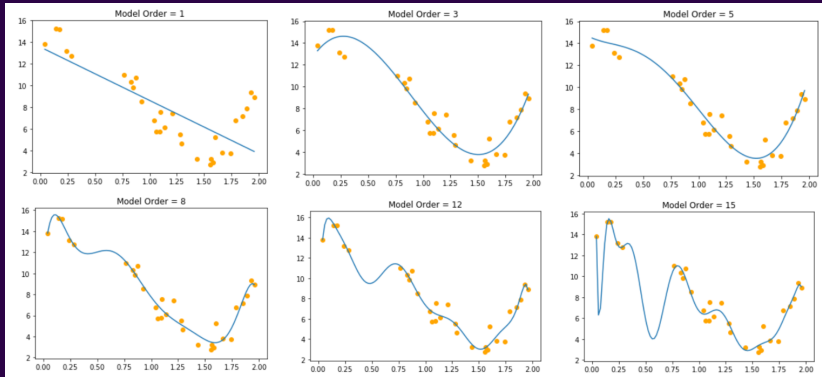
Outline

- 1 Review of Day 2
- 2 Lab: Robot Arm Calibration
- 3 Polynomial Regression
- 4 Train and Test Error, Overfitting
- 5 Model Order Selection
- 6 Regularization

Overfitting

- We learned how to fit our data using polynomials of different order
- With a higher model order, we can fit the data with increasing accuracy
- As you increase the model order, at certain point it is possible find a model that fits your data perfectly (ie. zero error)
- What could be the problem?


Overfitting



■ Which of these model do you think is the best? Why?


Overfitting

- The problem is that we are only fitting our model using data that is given
- Data usually contains noise
- When a model becomes too complex, it will start to fit the noise in the data
- What happens if we apply our model to predict some data that the model has never seen before? It will not work well.
- This is called over-fitting



Overfitting


- Split the data set into a train set and a test set
- Train set will be used to train the model
- The test set will not be seen by the model during the training process
- Use test set to evaluate the model when a model is trained



- With the training and test sets shown, which one do you think is the better model now?

Train and Test Error

- Plot of train error and test error for different model order
- Initially both train and test error go down as model order increase
- But at a certain point, test error start to increase because of overfitting



NYU

TANDON SCHOOL
OF ENGINEERING

Outline

- 1 Review of Day 2
- 2 Lab: Robot Arm Calibration
- 3 Polynomial Regression
- 4 Train and Test Error, Overfitting
- 5 Model Order Selection**
- 6 Regularization

Question:


- Can we **write an algorithm that automatically determines the correct model** order and uses this model?

Question:

- Can we **write an algorithm that automatically determines the correct model** order and uses this model?
- Test error often increases when we've passed the true model order ...


Question:

- Can we **write an algorithm that automatically determines the correct model** order and uses this model?
- Test error often increases when we've passed the true model order ...



Question:


- Can we **write an algorithm that automatically determines the correct model** order and uses this model?
- Test error often increases when we've passed the true model order ...



- Is optimizing our algorithm based on test error smart?

Question:


- Can we **write an algorithm that automatically determines the correct model** order and uses this model?
- Test error often increases when we've passed the true model order ...



- Is optimizing our algorithm based on test error smart?
 - We run into the **same problem as overfitting**

Question:

- Can we **write an algorithm that automatically determines the correct model** order and uses this model?
- Test error often increases when we've passed the true model order ...



- Is optimizing our algorithm based on test error smart?
 - We run into the **same problem as overfitting**
 - Tuning our algorithm on what should be unknown data!

Cross-Validation

- Motivation: never determine a hyper-parameter based on training data

Cross-Validation

- Motivation: never determine a hyper-parameter based on training data
- **Hyper-Parameter:** a parameter of the algorithm that is not a model-parameter solved for in optimization.

Cross-Validation

- Motivation: never determine a hyper-parameter based on training data
- **Hyper-Parameter:** a parameter of the algorithm that is not a model-parameter solved for in optimization.
 - Ex: model order vs. model weights (β)

Cross-Validation

- Motivation: never determine a hyper-parameter based on training data
- **Hyper-Parameter**: a parameter of the algorithm that is not a model-parameter solved for in optimization.
 - Ex: model order vs. model weights (β)
- Solution: split dataset into three

Cross-Validation

- Motivation: never determine a hyper-parameter based on training data
- **Hyper-Parameter**: a parameter of the algorithm that is not a model-parameter solved for in optimization.
 - Ex: model order vs. model weights (β)
- Solution: split dataset into three
 - **Training set**: to compute the model-parameters (β)

Cross-Validation

- Motivation: never determine a hyper-parameter based on training data
- **Hyper-Parameter**: a parameter of the algorithm that is not a model-parameter solved for in optimization.
 - Ex: model order vs. model weights (β)
- Solution: split dataset into three
 - **Training set**: to compute the model-parameters (β)
 - **Validation set**: to tune hyper-parameters (model-order)

Cross-Validation

- Motivation: never determine a hyper-parameter based on training data
- **Hyper-Parameter**: a parameter of the algorithm that is not a model-parameter solved for in optimization.
 - Ex: model order vs. model weights (β)
- Solution: split dataset into three
 - **Training set**: to compute the model-parameters (β)
 - **Validation set**: to tune hyper-parameters (model-order)
 - **Test set**: to compute the performance of the algorithm (MSE)

Cross-Validation

- Motivation: never determine a hyper-parameter based on training data
- **Hyper-Parameter**: a parameter of the algorithm that is not a model-parameter solved for in optimization.
 - Ex: model order vs. model weights (β)
- Solution: split dataset into three
 - **Training set**: to compute the model-parameters (β)
 - **Validation set**: to tune hyper-parameters (model-order)
 - **Test set**: to compute the performance of the algorithm (MSE)
- Demo: MOS Attempt 1

Finding a Rule for Model Order Selection

- Picking model-order with lowest val. MSE often over-predicts

Finding a Rule for Model Order Selection

- Picking model-order with lowest val. MSE often over-predicts
- How can we make our rule for MOS more reliable?

Finding a Rule for Model Order Selection

- Picking model-order with lowest val. MSE often over-predicts
- How can we make our rule for MOS more reliable?
- Are there any statistics we could use?

Finding a Rule for Model Order Selection

- Picking model-order with lowest val. MSE often over-predicts
- How can we make our rule for MOS more reliable?
- Are there any statistics we could use?
- **Possible Answer:** Fitting multiple datasets and averaging the validation error

K-Folds Cross-Validation Algorithm

- 1 Split Test+Val data into K sets, called **folds**

K-Folds Cross-Validation Algorithm

- 1 Split Test+Val data into K sets, called **folds**
- 2 For each Fold k

K-Folds Cross-Validation Algorithm

- 1 Split Test+Val data into K sets, called **folds**
- 2 For each Fold k
 - 1 Label Fold k as the validation set

K-Folds Cross-Validation Algorithm

- 1 Split Test+Val data into K sets, called **folds**
- 2 For each Fold k
 - 1 Label Fold k as the validation set
 - 2 Label every other fold as the training set

K-Folds Cross-Validation Algorithm

- 1 Split Test+Val data into K sets, called **folds**
- 2 For each Fold k
 - 1 Label Fold k as the validation set
 - 2 Label every other fold as the training set
 - 3 For each Model-Order m

K-Folds Cross-Validation Algorithm

- 1 Split Test+Val data into K sets, called **folds**
- 2 For each Fold k
 - 1 Label Fold k as the validation set
 - 2 Label every other fold as the training set
 - 3 For each Model-Order m
 - 1 Compute the model parameters (β)

K-Folds Cross-Validation Algorithm

- 1 Split Test+Val data into K sets, called **folds**
- 2 For each Fold k
 - 1 Label Fold k as the validation set
 - 2 Label every other fold as the training set
 - 3 For each Model-Order m
 - 1 Compute the model parameters (β)
 - 2 Compute the score on the val. set (MSE)

K-Folds Cross-Validation Algorithm

- 1 Split Test+Val data into K sets, called **folds**
- 2 For each Fold k
 - 1 Label Fold k as the validation set
 - 2 Label every other fold as the training set
 - 3 For each Model-Order m
 - 1 Compute the model parameters (β)
 - 2 Compute the score on the val. set (MSE)
- 3 Average the val. score over all folds, for each model-order

K-Folds Cross-Validation Algorithm

- 1 Split Test+Val data into K sets, called **folds**
- 2 For each Fold k
 - 1 Label Fold k as the validation set
 - 2 Label every other fold as the training set
 - 3 For each Model-Order m
 - 1 Compute the model parameters (β)
 - 2 Compute the score on the val. set (MSE)
- 3 Average the val. score over all folds, for each model-order
- 4 Find std-dev of lowest mean val. score

K-Folds Cross-Validation Algorithm

- 1 Split Test+Val data into K sets, called **folds**
- 2 For each Fold k
 - 1 Label Fold k as the validation set
 - 2 Label every other fold as the training set
 - 3 For each Model-Order m
 - 1 Compute the model parameters (β)
 - 2 Compute the score on the val. set (MSE)
- 3 Average the val. score over all folds, for each model-order
- 4 Find std-dev of lowest mean val. score
- 5 **Rule:** choose lowest model order with mean val. score within one SE of the lowest

K-Folds Cross-Validation Algorithm

- 1 Split Test+Val data into K sets, called **folds**
- 2 For each Fold k
 - 1 Label Fold k as the validation set
 - 2 Label every other fold as the training set
 - 3 For each Model-Order m
 - 1 Compute the model parameters (β)
 - 2 Compute the score on the val. set (MSE)
- 3 Average the val. score over all folds, for each model-order
- 4 Find std-dev of lowest mean val. score
- 5 **Rule:** choose lowest model order with mean val. score within one SE of the lowest
 - Standard-Error (SE):
 $(\text{std-dev of lowest mean val. score})/\sqrt{K-1}$

Outline

- 1 Review of Day 2
- 2 Lab: Robot Arm Calibration
- 3 Polynomial Regression
- 4 Train and Test Error, Overfitting
- 5 Model Order Selection
- 6 Regularization**

Can we prevent overfitting another way?

- **Regularization:** methods to prevent overfitting

Can we prevent overfitting another way?

- **Regularization:** methods to prevent overfitting
 - We just covered regularization by model order selection

Can we prevent overfitting another way?

- **Regularization:** methods to prevent overfitting
 - We just covered regularization by model order selection
- Running K-folds for cross-validation is intensive

Can we prevent overfitting another way?


- **Regularization:** methods to prevent overfitting
 - We just covered regularization by model order selection
- Running K-folds for cross-validation is intensive
- Is there another way? Talk among your classmates.

Can we prevent overfitting another way?

- **Regularization:** methods to prevent overfitting
 - We just covered regularization by model order selection
- Running K-folds for cross-validation is intensive
- Is there another way? Talk among your classmates.
 - Solution: We can change our cost function.


Weight Based Regularization

- Looking back at the polynomial overfitting



Weight Based Regularization

- Looking back at the polynomial overfitting
- Notice that weight-size increases with overfitting



New Cost Function

$$J = \sum_{i=1}^N (y_i - y_{i \text{ pred}})^2$$

New Cost Function

$$J = \sum_{i=1}^N (y_i - y_{i\text{pred}})^2 + \lambda \sum_{j=1}^D (w_j)^2$$

New Cost Function

$$J = \sum_{i=1}^N (y_i - y_{i\text{pred}})^2 + \lambda \sum_{j=1}^D (w_j)^2$$

- Penalize complexity by simultaneously minimizing weight values.

New Cost Function

$$J = \sum_{i=1}^N (y_i - y_{i\text{pred}})^2 + \lambda \sum_{j=1}^D (w_j)^2$$

- Penalize complexity by simultaneously minimizing weight values.
- We call λ a **hyperparameter**

New Cost Function

$$J = \sum_{i=1}^N (y_i - y_{i_{pred}})^2 + \lambda \sum_{j=1}^D (w_j)^2$$

- Penalize complexity by simultaneously minimizing weight values.
- We call λ a **hyperparameter**
 - λ determines relative importance

New Cost Function

$$J = \sum_{i=1}^N (y_i - y_{i_{pred}})^2 + \lambda \sum_{j=1}^D (w_j)^2$$

- Penalize complexity by simultaneously minimizing weight values.
- We call λ a **hyperparameter**
 - λ determines relative importance

New Cost Function

$$J = \sum_{i=1}^N (y_i - y_{i\text{pred}})^2 + \lambda \sum_{j=1}^D (w_j)^2$$

- Penalize complexity by simultaneously minimizing weight values.
- We call λ a **hyperparameter**
 - λ determines relative importance

Table 1.2 Table of the coefficients w^* for $M = 9$ polynomials with various values for the regularization parameter λ . Note that $\ln \lambda = -\infty$ corresponds to a model with no regularization, i.e., to the graph at the bottom right in Figure 1.4. We see that, as the value of λ increases, the typical magnitude of the coefficients gets smaller.

	$\ln \lambda = -\infty$	$\ln \lambda = -18$	$\ln \lambda = 0$
w_0^*	0.35	0.35	0.13
w_1^*	232.37	4.74	-0.05
w_2^*	-5321.83	-0.77	-0.06
w_3^*	48568.31	-31.97	-0.05
w_4^*	-231639.30	-3.89	-0.03
w_5^*	640042.26	55.28	-0.02
w_6^*	-1061800.52	41.32	-0.01
w_7^*	1042400.18	-45.95	-0.00
w_8^*	-557682.99	-91.53	0.00
w_9^*	125201.43	72.68	0.01



Learning Objectives

- What is the difference between train error and test error?
- What is overfitting? How do we detect it?
- What is cross validation?
- How to find the optimal model order for my model?
- What is regularization? How does it prevent overfitting?

Thank You!

- Next Class: Linear Classification