

# Day 4: Linear Classifiers

## Summer STEM: Machine Learning

Nikola Janjušević, Akshaj Kumar Veldanda, Jacky Yuan,  
Tejaishwarya Gagadam

Department of Electrical and Computer Engineering  
NYU Tandon School of Engineering  
Brooklyn, New York

July 11, 2019

# Outline

- 1 Review of Day 3
- 2 Non-Linear Optimization
- 3 Demo: Diagnosing Breast Cancer
- 4 Logistic Regression
- 5 Decision Thresholds and ROC
- 6 Lab: Titanic
- 7 Multiclass Classification
- 8 Lab: Multiclass

# Review of Day 3

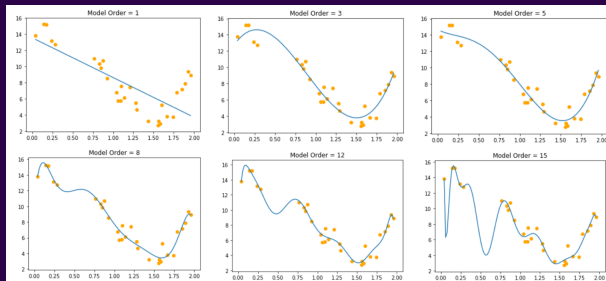
- Yesterday we learned about:
  -
- Polynomial Regression
- Overfitting
- Regularization
- Cross-Validation

# Review: Polynomial Regression

- Polynomial Model:  $y = w_0 + w_1x + w_2x^2 + w_3x^3 + \dots w_Dx^D$
- Design Matrix for Polynomial:  $X = \begin{bmatrix} 1 & x_1 & x_1^2 & \dots & x_1^D \\ 1 & x_2 & x_2^2 & \dots & x_2^D \\ \vdots & & \ddots & & \vdots \\ 1 & x_N & x_N^2 & \dots & x_N^D \end{bmatrix}$
- Think of this design matrix as creating new features that are powers of the original single feature
- This is Linear Regression, the process for calculating the weights  $w$ 's are same
- Python function: `sklearn` or `np.polyfit`

# Review: Over-fitting

- Train error always decreases as you use higher order models
- A model that is over-fitted on train data is unlikely to work well on new data



# Cross-Validation

- Motivation: never determine a hyper-parameter based on training data
- **Hyper-Parameter**: a parameter of the algorithm that is not a model-parameter solved for in optimization.
  - Ex:  $\lambda$  weight regularization value vs. model weights ( $\mathbf{w}$ )
- Solution: split dataset into three
  - **Training set**: to compute the model-parameters ( $\mathbf{w}$ )
  - **Validation set**: to tune hyper-parameters ( $\lambda$ )
  - **Test set**: to compute the performance of the algorithm (MSE)

# Review: Regularization

- Another method to combat over-fitting
- High weight terms usually lead to over-fitting
- Introduction of a new term in cost function
- 

$$J = \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

# Review: Regularization

- Another method to combat over-fitting
- High weight terms usually lead to over-fitting
- Introduction of a new term in cost function
- 

$$J = \sum_{i=1}^N (y_i - \hat{y}_i)^2 + \lambda \sum_{j=1}^D (w_j)^2$$

- The new term penalizes the magnitude of weights
- Hyperparameter *lambda* determines how much you regularize:  
higher lambda  $\leftarrow$  more regularization



# Outline

- 1 Review of Day 3
- 2 **Non-Linear Optimization**
- 3 Demo: Diagnosing Breast Cancer
- 4 Logistic Regression
- 5 Decision Thresholds and ROC
- 6 Lab: Titanic
- 7 Multiclass Classification
- 8 Lab: Multiclass

# Motivation

- Cannot rely on closed form solutions

# Motivation

- Cannot rely on closed form solutions
  - Computation efficiency: operations like inverting a matrix is not efficient

# Motivation

- Cannot rely on closed form solutions
  - Computation efficiency: operations like inverting a matrix is not efficient
  - For more complex problems, like neural network, a closed form solution is not always available

# Motivation



- Cannot rely on closed form solutions
  - Computation efficiency: operations like inverting a matrix is not efficient
  - For more complex problems, like neural network, a closed form solution is not always available
- Need an optimization technique to find an optimal solution

# Motivation

- Cannot rely on closed form solutions
  - Computation efficiency: operations like inverting a matrix is not efficient
  - For more complex problems, like neural network, a closed form solution is not always available
- Need an optimization technique to find an optimal solution
  - Machine learning practitioners use gradient based methods

# Understanding Optimization

- *Recap*  $\hat{y} = w_0 + w_1x$
- *Loss*,  $J = \sum_{i=1}^N (y_i - \hat{y}_i)^2 \implies J = \sum_{i=1}^N (y_i - w_0 - w_1x_i)^2$
- Want to find  $w_0$  and  $w_1$  that minimizes  $J$



# Gradient Descent Algorithm


## ■ Update Rule

*Repeat*{

$$w_{new} = w - \alpha \frac{dJ}{dw}$$

}


$\alpha$  is the learning rate







# General Loss Function Contours

- Most loss function contours are not perfectly parabolic
- Our goal is to find a solution that is very close to global minimum by the right choice of hyper parameters



# Understanding Learning Rate



# Some Animations

- Demonstrate gradient descent animation


# Importance of Feature Normalization (Optional)

- Helps improve the performance of gradient based optimization

Without feature scaling



With feature scaling



# Some Gradient Based Algorithms

- Gradient descent
  - Stochastic gradient descent
  - Mini-batch gradient descent
- Gradient descent with momentum
- RMSprop
- Adam optimization algorithm

We have many frameworks that help us use these techniques in a single line of code (Eg: TensorFlow, PyTorch, Caffe, etc).

# Outline

- 1 Review of Day 3
- 2 Non-Linear Optimization
- 3 Demo: Diagnosing Breast Cancer**
- 4 Logistic Regression
- 5 Decision Thresholds and ROC
- 6 Lab: Titanic
- 7 Multiclass Classification
- 8 Lab: Multiclass

# Demo: Diagnosing Breast Cancer

- We're going to use the breast cancer dataset to predict whether the patients' scans show a malignant tumour or a benign tumour.
- Let's try to find the best linear classifier using linear regression.

# Outline

- 1 Review of Day 3
- 2 Non-Linear Optimization
- 3 Demo: Diagnosing Breast Cancer
- 4 Logistic Regression**
- 5 Decision Thresholds and ROC
- 6 Lab: Titanic
- 7 Multiclass Classification
- 8 Lab: Multiclass



# Classification

- One method is to use linear regression
  - Map all predictions ( $\hat{y}$ ) greater than 0 as class 1 and all less than 0 as class 0
  - This method doesn't work well because classification is not actually a linear function
- Classification takes the only discrete values for prediction
  - For binary classification problem,  $y$  can only take two values, 0 and 1
  - Ex: If we want to build a spam classifier for email, then  $x$  may be some features of the email
    - The  $y = 1$  if it is a spam
    - Otherwise,  $y = 0$

# Hypothesis Representation

- Approach classification as old linear regression problem, ignoring the fact that  $y$  is discrete

# Hypothesis Representation

- Approach classification as old linear regression problem, ignoring the fact that  $y$  is discrete
  - We have seen that this approach performs poorly

# Hypothesis Representation



- Approach classification as old linear regression problem, ignoring the fact that  $y$  is discrete
  - We have seen that this approach performs poorly
- To fix this, develop an hypothesis such that  $0 \leq \hat{y} \leq 1$

# Hypothesis Representation

- Approach classification as old linear regression problem, ignoring the fact that  $y$  is discrete
  - We have seen that this approach performs poorly
- To fix this, develop an hypothesis such that  $0 \leq \hat{y} \leq 1$ 
  - This is accomplished by using the Sigmoid function



# Sigmoid Function

- Recall from linear regression  $z = w_0 + w_1x$
- On application of sigmoid function to  $z$ , we force  $0 \leq \hat{y} \leq 1$ 
  - $\hat{y} = \text{sigmoid}(z) = \frac{1}{1+e^{-z}}$



# Classification Loss Function

- Cannot use the same cost function that we used for linear regression
  - Logistic function has many local optima
- Logistic cost function is  $\frac{1}{m} \sum_{i=1}^m [-y \log(\hat{y}) - (1 - y) \log(1 - \hat{y})]$ 
  - This loss function is called binary cross entropy loss
  - This loss function has only one optimum point



# Outline

- 1 Review of Day 3
- 2 Non-Linear Optimization
- 3 Demo: Diagnosing Breast Cancer
- 4 Logistic Regression
- 5 Decision Thresholds and ROC**
- 6 Lab: Titanic
- 7 Multiclass Classification
- 8 Lab: Multiclass



# Types of Errors in Classification

- Correct predictions:
  - True Positive (TP) : Predict  $y = 1$  when  $y = 1$
  - True Negative (TN) : Predict  $y = 0$  when  $y = 0$
- Two types of errors:
  - False Positive/False Alarm (FP): Predict  $y=1$  when  $y=0$
  - False Negative/Missed Detection (FN): Predict  $y=0$  when  $y=1$
- Confusion Matrix:

□ Accuracy of classifier can be measured by:

- $TPR = P(\hat{y} = 1|y = 1)$
- $FPR = P(\hat{y} = 1|y = 0)$
- $Accuracy = P(\hat{y} = 1|y = 1) + P(\hat{y} = 0|y = 0)$ 
  - (percentage of correct classification)

predicted→ real↓	Class_pos	Class_neg
Class_pos	TP	FN
Class_neg	FP	TN

$$TPR \text{ (sensitivity)} = \frac{TP}{TP + FN}$$

$$FPR \text{ (1-specificity)} = \frac{FP}{TN + FP}$$

# Different Metrics for Error


- Metrics to measure the error rate:
  - Recall/Sensitivity/TPR =  $TP/(TP+FN)$  (How many positives are detected among all positive?)
  - Precision =  $TP/(TP+FP)$  (How many detected positive is actually positive?)
  - Accuracy =  $(TP+TF)/(TP+FP+TN+FN)$  (percentage of correct classification)
  - F1-score =  $\frac{Precision * Recall}{(Precision + Recall)/2}$
- Why accuracy alone is not a good measure for assessing the model
  - There might be an overwhelming proportion of one class over another
  - Example: A rare disease occurs 1 in ten thousand people
  - A test that classifies everyone as free of the disease can achieve 99.999% accuracy when tested with people drawn randomly from the entire population

# Thresholding and ROC

- We can trade-off TPR (sensitivity) and FPR by changing the threshold
- Increasing  $t \rightarrow$  Decreases false positives, but also reduces sensitivity
- Decreasing  $t \rightarrow$  Increases sensitivity, but also increases false positive
- Why do we want this trade-off?
- Example:
  - 1 Detection for burglary into a building, need high sensitivity, we can tolerate a few false alarms  $\rightarrow$  decrease  $t$
  - 2 Making decision to buy a stock, making a false positive decision will lose millions  $\rightarrow$  increase  $t$

# Thresholding and ROC

- ROC (Receiver Operating Characteristics) curve:
- Plot the change between TPR and FPR by varying the threshold
- Allow you to choose the threshold to meet a target TPR/FPR
- A good classifier will have large area under the curve
- A classifier with a higher area under the curve means that under same FPR, it has higher TPR



# Outline

- 1 Review of Day 3
- 2 Non-Linear Optimization
- 3 Demo: Diagnosing Breast Cancer
- 4 Logistic Regression
- 5 Decision Thresholds and ROC
- 6 Lab: Titanic**
- 7 Multiclass Classification
- 8 Lab: Multiclass

# Outline

- 1 Review of Day 3
- 2 Non-Linear Optimization
- 3 Demo: Diagnosing Breast Cancer
- 4 Logistic Regression
- 5 Decision Thresholds and ROC
- 6 Lab: Titanic
- 7 Multiclass Classification**
- 8 Lab: Multiclass

# On Board:

- Previous Model:  $\hat{y} = \sigma(\mathbf{x}^T \mathbf{w})$

# On Board:

- Previous Model:  $\hat{y} = \sigma(\mathbf{x}^T \mathbf{w})$
- Representing Multiple Classes:



# On Board:

- Previous Model:  $\hat{y} = \sigma(\mathbf{x}^T \mathbf{w})$
- Representing Multiple Classes:
  - One-hot / 1-of-K vectors, ex: Class 2 of 4  $[0, 1, 0, 0]$

# On Board:

- Previous Model:  $\hat{y} = \sigma(\mathbf{x}^T \mathbf{w})$
- Representing Multiple Classes:
  - One-hot / 1-of-K vectors, ex: Class 2 of 4  $[0, 1, 0, 0]$
- Multiple Outputs:

$$\hat{\mathbf{y}} = \text{softmax}(W\mathbf{x})$$

# On Board:

- Previous Model:  $\hat{y} = \sigma(\mathbf{x}^T \mathbf{w})$
- Representing Multiple Classes:
  - One-hot / 1-of-K vectors, ex: Class 2 of 4  $[0, 1, 0, 0]$
- Multiple Outputs:

$$\hat{\mathbf{y}} = \text{softmax}(W\mathbf{x})$$

- $(K,1) = (K,M) \times (M,1)$

# On Board:

- Previous Model:  $\hat{y} = \sigma(\mathbf{x}^T \mathbf{w})$
- Representing Multiple Classes:
  - One-hot / 1-of-K vectors, ex: Class 2 of 4  $[0, 1, 0, 0]$
- Multiple Outputs:

$$\hat{\mathbf{y}} = \text{softmax}(\mathbf{W}\mathbf{x})$$

- $(K,1) = (K,M) \times (M,1)$

$$\text{softmax}(\mathbf{z})_k = \frac{e^{z_k}}{\sum_j e^{z_j}}$$

# On Board:

- Previous Model:  $\hat{y} = \sigma(\mathbf{x}^T \mathbf{w})$
- Representing Multiple Classes:
  - One-hot / 1-of-K vectors, ex: Class 2 of 4  $[0, 1, 0, 0]$
- Multiple Outputs:

$$\hat{\mathbf{y}} = \text{softmax}(\mathbf{W}\mathbf{x})$$

- $(K,1) = (K,M) \times (M,1)$

$$\text{softmax}(\mathbf{z})_k = \frac{e^{z_k}}{\sum_j e^{z_j}}$$

- Cross-Entropy:  $J = -\sum_{i=1}^N \sum_{k=1}^K y_{ik} \ln(\hat{y}_{ik})$

# Outline

- 1 Review of Day 3
- 2 Non-Linear Optimization
- 3 Demo: Diagnosing Breast Cancer
- 4 Logistic Regression
- 5 Decision Thresholds and ROC
- 6 Lab: Titanic
- 7 Multiclass Classification
- 8 Lab: Multiclass**

# Thank You!

- Next Class: Linear Regression
- The real machine learning will begin!