

Day 4: Linear Classifiers

Summer STEM: Machine Learning

Nikola Janjušević
Akshaj Kumar Veldanda
Jacky Yuan
Tejaishwarya Gagadam

NYU Tandon School of Engineering

June 20, 2019

Learning Objectives

- What is the functional form of logistic regression?
- How do we interpret the output of a binary logistic classifier? Multi-variable classifier?
- How do we determine the threshold for classification?
- What is cross-entropy loss function? Why do we use cross-entropy over MSE?
- What is one-hot vector?
- How to do multi-class classification?
- What is an ROC curve and how do we interpret it?

Outline

- 1 Review of Day 3
- 2 Demo: Diagnosing Breast Cancer
- 3 Logistic Regression
- 4 Decision Thresholds and ROC
- 5 Lab: *Singleclass*
- 6 Multiclass Classification
- 7 Lab: *multiclass*

Review of Day 3

- Yesterday we learned about:
- Polynomial Regression
- Overfitting
- K-folds Validation
- Regularization
- Non-Linear Optimization - Gradient Descent

Review: Polynomial Regression

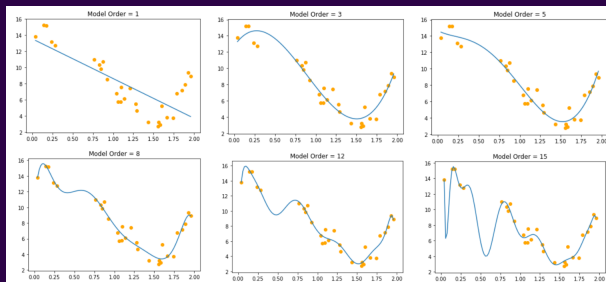
- Polynomial Model: $y = \beta_0 + \beta_1 x + \beta_2 x^2 + \beta_3 x^3 + \dots \beta_D x^D$

- Design Matrix for Polynomial: $A = \begin{bmatrix} 1 & x_1 & x_1^2 & \dots & x_1^D \\ 1 & x_2 & x_2^2 & \dots & x_2^D \\ \vdots & & \ddots & & \vdots \\ 1 & x_N & x_N^2 & \dots & x_N^D \end{bmatrix}$

- Think of this design matrix as creating new features that are powers of the original single feature
- The process for calculating the weights β 's are same as Linear Regression
- Python function: `sklearn` or `np.polyfit`

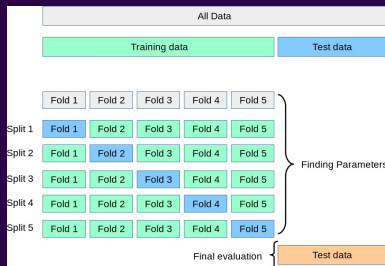
Review: Over-fitting

- Train error always decreases as you use higher order models
- A model that is over-fitted on train data is unlikely to work well on new data



Review: K-folds Validation

- Algorithm to automatically detect the optimal model order
- Split train data into K folds
- On each iteration, evaluate test error using one of the folds
- Calculate the average error among all folds



Review: Regularization

- Another method to combat over-fitting
- High weight terms usually lead to over-fitting
- Introduction of a new term in cost function
-

$$J = \sum_{i=1}^N (y_i - y_{i\text{pred}})^2$$

Review: Regularization

- Another method to combat over-fitting
- High weight terms usually lead to over-fitting
- Introduction of a new term in cost function
-

$$J = \sum_{i=1}^N (y_i - y_{i\text{pred}})^2 + \lambda \sum_{j=1}^D (w_j)^2$$

- The new term penalizes the magnitude of weights
- Hyperparameter *lambda* determines how much you regularize:
higher lambda \leftarrow more regularization

Review: Non-Linear Optimization


- Cannot rely on closed form solutions
 - Computation efficiency
 - Closed form solution not available
- Optimization technique finds optimal solution by iteratively changing parameters toward better solution
- Gradient based method is most commonly used in ML

Review: Gradient Descent

- Average rate of change: $\frac{\Delta y}{\Delta x} = \frac{y_1 - y_0}{x_1 - x_0}$
- Derivative: Instantaneous rate of change of a function
- $\frac{dy}{dx} = \lim_{x \rightarrow 0} \frac{f(x + \Delta x) - f(x)}{\Delta x}$
- For a 1-dimensional function:
 - $\frac{dy}{dx} > 0$, Function is increasing
 - $\frac{dy}{dx} < 0$, Function is decreasing


Review: Gradient Descent

- In a two dimensional function $z = f(x, y)$, we can take derivative in more than one direction
- Gradient a vector that points to the direction of maximum increase
- $\nabla f = \left[\frac{\partial f}{\partial x} \quad \frac{\partial f}{\partial y} \right]^T$



Review: Gradient Descent

- Weight Space
- We may visualize the loss function as surface in a multi-dimensional space
- Locally, the function may be viewed as a paraboloid
- There are local minima that we want to avoid because they are not optimal solution



Review: Gradient Descent


- Key idea: For each iteration

- 1. Calculate the gradient of the loss function with respect to the weights

$$\nabla J(\mathbf{w}) = \left[\frac{\partial J}{\partial \beta_0} \quad \cdots \quad \frac{\partial J}{\partial \beta_N} \right]^T$$

- 2. Update the weights by taking a step opposite to the gradient $\mathbf{w}_{new} = \mathbf{w}_{old} - \alpha \nabla J(\mathbf{w})$
where *alpha* is called the learning rate: how big of a step you take

Review: Gradient Descent



Outline

- 1 Review of Day 3
- 2 Demo: Diagnosing Breast Cancer
- 3 Logistic Regression
- 4 Decision Thresholds and ROC
- 5 Lab: *Singleclass*
- 6 Multiclass Classification
- 7 Lab: *multiclass*

Demo: Diagnosing Breast Cancer

- We're going to use the breast cancer dataset to predict whether the patients' scans show a malignant tumour or a benign tumour.
- Let's try to find the best linear classifier using linear regression.

Outline

- 1 Review of Day 3
- 2 Demo: Diagnosing Breast Cancer
- 3 Logistic Regression
- 4 Decision Thresholds and ROC
- 5 Lab: *Singleclass*
- 6 Multiclass Classification
- 7 Lab: *multiclass*

Classification



- One method is to use linear regression
 - Map all predictions greater than 0.5 as a 1 and all less than 0.5 as a 0
 - this method doesn't work well because classification is not actually a linear function
- Classification takes the only discrete values for prediction
 - For binary classification problem, y can only take two values, 0 and 1
 - Ex: If we want to build a spam classifier for email, then x may be some features of the email
 - The $y = 1$ if it is a spam
 - Otherwise, $y = 0$

Hypothesis Representation

- Approach classification as old linear regression problem, ignoring the fact that y is discrete
 - We have seen that this approach performs poorly
- To fix this, develop an hypothesis such that $0 \leq \hat{y} \leq 1$
 - This is accomplished by using the Sigmoid function


Sigmoid Function

- Recall from linear regression $z = \beta_0 + \beta_1 x$
- On application of sigmoid function to z , we force $0 \leq \hat{y} \leq 1$
 - $\hat{y} = \text{sigmoid}(z) = \frac{1}{1+e^{-z}}$



Classification Loss Function

- Cannot use the same cost function that we used for linear regression
 - Logistic function has many local optima
- Logistic cost function is $\frac{1}{m} \sum_{i=1}^m [-y \log(\hat{y}) - (1 - y) \log(1 - \hat{y})]$
 - This loss function is called binary cross entropy loss
 - This loss function has only one optimum point



Outline

- 1 Review of Day 3
- 2 Demo: Diagnosing Breast Cancer
- 3 Logistic Regression
- 4 Decision Thresholds and ROC
- 5 Lab: *Singleclass*
- 6 Multiclass Classification
- 7 Lab: *multiclass*

Types of Errors in Classification

- Correct predictions:
 - True Positive (TP) : Predict $y = 1$ when $y = 1$
 - True Negative (TN) : Predict $y = 0$ when $y = 0$
- Two types of errors:
 - False Positive/False Alarm (FP): Predict $y=1$ when $y=0$
 - False Negative/Missed Detection (FN): Predict $y=0$ when $y=1$
- Confusion Matrix:

□ Accuracy of classifier can be measured by:

- $TPR = P(\hat{y} = 1|y = 1)$
- $FPR = P(\hat{y} = 1|y = 0)$
- $Accuracy = P(\hat{y} = 1|y = 1) + P(\hat{y} = 0|y = 0)$
 - (percentage of correct classification)

predicted real → ↓	Class_pos	Class_neg
Class_pos	TP	FN
Class_neg	FP	TN

$$TPR \text{ (sensitivity)} = \frac{TP}{TP + FN}$$

$$FPR \text{ (1-specificity)} = \frac{FP}{TN + FP}$$

Different Metrics for Error


- Metrics to measure the error rate:
 - Recall/Sensitivity/TPR = $TP/(TP+FN)$ (How many positives are detected among all positive?)
 - Precision = $TP/(TP+FP)$ (How many detected positive is actually positive?)
 - Accuracy = $(TP+TF)/(TP+FP+TN+FN)$ (percentage of correct classification)
 - F1-score = $\frac{Precision * Recall}{(Precision + Recall)/2}$
- Why accuracy alone is not a good measure for assessing the model
 - There might be an overwhelming proportion of one class over another
 - Example: A rare disease occurs 1 in ten thousand people
 - A test that classifies everyone as free of the disease can achieve 99.999% accuracy when tested with people drawn randomly from the entire population

Thresholding and ROC

- We can trade-off TPR (sensitivity) and FPR by changing the threshold
- Increasing $t \rightarrow$ Decreases false positives, but also reduces sensitivity
- Decreasing $t \rightarrow$ Increases sensitivity, but also increases false positive
- Why do we want this trade-off?
- Example:
 - 1 Detection for burglary into a building, need high sensitivity, we can tolerate a few false alarms \rightarrow decrease t
 - 2 Making decision to buy a stock, making a false positive decision will lose millions \rightarrow increase t

Thresholding and ROC

- ROC (Receiver Operating Characteristics) curve:
- Plot the change between TPR and FPR by varying the threshold
- Allow you to choose the threshold to meet a target TPR/FPR
- A good classifier will have large area under the curve
- A classifier with a higher area under the curve means that under same FPR, it has higher TPR



Outline

- 1 Review of Day 3
- 2 Demo: Diagnosing Breast Cancer
- 3 Logistic Regression
- 4 Decision Thresholds and ROC
- 5 Lab: *Singleclass***
- 6 Multiclass Classification
- 7 Lab: *multiclass*

Outline

- 1 Review of Day 3
- 2 Demo: Diagnosing Breast Cancer
- 3 Logistic Regression
- 4 Decision Thresholds and ROC
- 5 Lab: *Singleclass*
- 6 Multiclass Classification
- 7 Lab: *multiclass*

Outline

- 1 Review of Day 3
- 2 Demo: Diagnosing Breast Cancer
- 3 Logistic Regression
- 4 Decision Thresholds and ROC
- 5 Lab: *Singleclass*
- 6 Multiclass Classification
- 7 Lab: *multiclass*

Learning Objectives

- What is the functional form of logistic regression?
- How do we interpret the output of a binary logistic classifier? Multi-variable classifier?
- How do we determine the threshold for classification?
- What is cross-entropy loss function? Why do we use cross-entropy over MSE?
- What is one-hot vector?
- How to do multi-class classification?
- What is an ROC curve and how do we interpret it?

