

Day 6: Neural Networks

Summer STEM: Machine Learning

Nikola Janjušević, Akshaj Kumar Veldanda, Jacky Yuan,
Tejaishwarya Gagadam

Department of Electrical and Computer Engineering
NYU Tandon School of Engineering
Brooklyn, New York

June 24, 2019

Learning Objectives

- What are the advantages of Neural Networks?
- What is the mathematical model for a Neural Network?
- What are the hyper-parameters associated with a Neural Network?
- Why are batch-size and learning-rate important? How are they related?
- How do we implement Neural Networks with Keras?

Outline

- 1 Review of Week 1
- 2 Neural Network Model
- 3 Training with Neural Networks
- 4 Introduction to Keras
- 5 Lab: Music Classification
- 6 (Optional) Lab: Cat vs. Non-Cat

Regression or Classification?

- **Problem 1:** Categorizing credit card applications into those who have good credit, bad credit and those who fall in the gray area.

Regression or Classification?

- **Problem 1:** Categorizing credit card applications into those who have good credit, bad credit and those who fall in the gray area.
 - Classification Problem

Regression or Classification?

- **Problem 1:** Categorizing credit card applications into those who have good credit, bad credit and those who fall in the gray area.
 - Classification Problem
 - Class Labels: Good credit, Bad credit, Average credit

Regression or Classification?

- **Problem 1:** Categorizing credit card applications into those who have good credit, bad credit and those who fall in the gray area.
 - Classification Problem
 - Class Labels: Good credit, Bad credit, Average credit
- **Problem 2:** Determining the sentiment of customer reviews for a product on Amazon.

Regression or Classification?

- **Problem 1:** Categorizing credit card applications into those who have good credit, bad credit and those who fall in the gray area.
 - Classification Problem
 - Class Labels: Good credit, Bad credit, Average credit
- **Problem 2:** Determining the sentiment of customer reviews for a product on Amazon.
 - Classification Problem

Regression or Classification?

- **Problem 1:** Categorizing credit card applications into those who have good credit, bad credit and those who fall in the gray area.
 - Classification Problem
 - Class Labels: Good credit, Bad credit, Average credit
- **Problem 2:** Determining the sentiment of customer reviews for a product on Amazon.
 - Classification Problem
 - Class Labels: Positive, Negative, Neutral

Regression or Classification?

- **Problem 1:** Categorizing credit card applications into those who have good credit, bad credit and those who fall in the gray area.
 - Classification Problem
 - Class Labels: Good credit, Bad credit, Average credit
- **Problem 2:** Determining the sentiment of customer reviews for a product on Amazon.
 - Classification Problem
 - Class Labels: Positive, Negative, Neutral
- **Problem 3:** Predict whether an employee's income is over 100k a year or not.

Regression or Classification?

- **Problem 1:** Categorizing credit card applications into those who have good credit, bad credit and those who fall in the gray area.
 - Classification Problem
 - Class Labels: Good credit, Bad credit, Average credit
- **Problem 2:** Determining the sentiment of customer reviews for a product on Amazon.
 - Classification Problem
 - Class Labels: Positive, Negative, Neutral
- **Problem 3:** Predict whether an employee's income is over 100k a year or not.
 - Classification Problem

Regression or Classification?

- **Problem 1:** Categorizing credit card applications into those who have good credit, bad credit and those who fall in the gray area.
 - Classification Problem
 - Class Labels: Good credit, Bad credit, Average credit
- **Problem 2:** Determining the sentiment of customer reviews for a product on Amazon.
 - Classification Problem
 - Class Labels: Positive, Negative, Neutral
- **Problem 3:** Predict whether an employee's income is over 100k a year or not.
 - Classification Problem
 - Class Labels: Over 100k, Under 100k

Regression or Classification?

- **Problem 4:** Estimating change in climate.

Regression or Classification?

- **Problem 4:** Estimating change in climate.
 - Regression Problem

Regression or Classification?

- **Problem 4:** Estimating change in climate.
 - Regression Problem
 - Target Variable: Predicting future temperatures

Regression or Classification?

- **Problem 4:** Estimating change in climate.
 - Regression Problem
 - Target Variable: Predicting future temperatures
- **Problem 5:** Identifying hate speech in social media.

Regression or Classification?

- **Problem 4:** Estimating change in climate.
 - Regression Problem
 - Target Variable: Predicting future temperatures
- **Problem 5:** Identifying hate speech in social media.
 - Classification Problem

Regression or Classification?

- **Problem 4:** Estimating change in climate.
 - Regression Problem
 - Target Variable: Predicting future temperatures
- **Problem 5:** Identifying hate speech in social media.
 - Classification Problem
 - Class labels: Normal Speech, Hate Speech

Regression or Classification?

- **Problem 4:** Estimating change in climate.
 - Regression Problem
 - Target Variable: Predicting future temperatures
- **Problem 5:** Identifying hate speech in social media.
 - Classification Problem
 - Class labels: Normal Speech, Hate Speech
- **Problem 6:** Forecasting the energy demand in a region.

Regression or Classification?

- **Problem 4:** Estimating change in climate.
 - Regression Problem
 - Target Variable: Predicting future temperatures
- **Problem 5:** Identifying hate speech in social media.
 - Classification Problem
 - Class labels: Normal Speech, Hate Speech
- **Problem 6:** Forecasting the energy demand in a region.
 - Regression Problem

Regression or Classification?

- **Problem 4:** Estimating change in climate.
 - Regression Problem
 - Target Variable: Predicting future temperatures
- **Problem 5:** Identifying hate speech in social media.
 - Classification Problem
 - Class labels: Normal Speech, Hate Speech
- **Problem 6:** Forecasting the energy demand in a region.
 - Regression Problem
 - Target Variable: Predicting the amount of energy needed in the future

Machine Learning Problem Pipeline

- 1 Gather data
- 2 Visualize the data
- 3 Formulate ML problem
 - Regression vs Classification
 - Choose an appropriate cost function
- 4 Design the model and train to find the optimal parameters of the model
 - Prepare a design matrix
 - Perform feature engineering
 - Validate your choice of hyper-parameters using a cross-validation set
- 5 Evaluate the model on a test set
 - If the performance is not satisfactory, go back to step 4

Data

- Always save your data file as an .csv file
 - It is easy to edit in both excel and text file
 - Easy to load the data using Pandas
- Visualize the data
 - To get an rough estimate of how your machine learning model should be
 - Do you have sufficient training and testing data
- Always plot the data before pre-processing

Notation

■ Numbers:

Notation

- Numbers:
 - N : total number of samples

Notation

- Numbers:
 - N : total number of samples
 - M : model order, number of features (engineered or not)

Notation

- Numbers:
 - N : total number of samples
 - M : model order, number of features (engineered or not)
 - K : number of outputs or classes

Notation

- Numbers:
 - N : total number of samples
 - M : model order, number of features (engineered or not)
 - K : number of outputs or classes
- Vectors:

Notation

- Numbers:
 - N: total number of samples
 - M: model order, number of features (engineered or not)
 - K: number of outputs or classes
- Vectors:
 - \mathbf{x} : feature vector, $\mathbf{x} = [x_1, x_2, \dots, x_M]^T$

Notation

■ Numbers:

- N: total number of samples
- M: model order, number of features (engineered or not)
- K: number of outputs or classes

■ Vectors:

- \mathbf{x} : feature vector, $\mathbf{x} = [x_1, x_2, \dots, x_M]^T$
- \mathbf{y} : target vector, $\mathbf{y} = [y_1, y_2, \dots, y_K]^T$

Notation

■ Numbers:

- N: total number of samples
- M: model order, number of features (engineered or not)
- K: number of outputs or classes

■ Vectors:

- \mathbf{x} : feature vector, $\mathbf{x} = [x_1, x_2, \dots, x_M]^T$
- \mathbf{y} : target vector, $\mathbf{y} = [y_1, y_2, \dots, y_K]^T$
- $\hat{\mathbf{y}}$: predicted target vector, $\hat{\mathbf{y}} = [\hat{y}_1, \hat{y}_2, \dots, \hat{y}_K]^T$

Notation

■ Numbers:

- N : total number of samples
- M : model order, number of features (engineered or not)
- K : number of outputs or classes

■ Vectors:

- \mathbf{x} : feature vector, $\mathbf{x} = [x_1, x_2, \dots, x_M]^T$
- \mathbf{y} : target vector, $\mathbf{y} = [y_1, y_2, \dots, y_K]^T$
- $\hat{\mathbf{y}}$: predicted target vector, $\hat{\mathbf{y}} = [\hat{y}_1, \hat{y}_2, \dots, \hat{y}_K]^T$
- \mathbf{w} : weight vector for $K = 1$ targets, $\mathbf{w} = [w_1, w_2, \dots, w_M]^T$

Notation

■ Numbers:

- N: total number of samples
- M: model order, number of features (engineered or not)
- K: number of outputs or classes

■ Vectors:

- \mathbf{x} : feature vector, $\mathbf{x} = [x_1, x_2, \dots, x_M]^T$
- \mathbf{y} : target vector, $\mathbf{y} = [y_1, y_2, \dots, y_K]^T$
- $\hat{\mathbf{y}}$: predicted target vector, $\hat{\mathbf{y}} = [\hat{y}_1, \hat{y}_2, \dots, \hat{y}_K]^T$
- \mathbf{w} : weight vector for $K = 1$ targets, $\mathbf{w} = [w_1, w_2, \dots, w_M]^T$
 - Also used: $\boldsymbol{\beta} = [\beta_1, \beta_2, \dots, \beta_M]^T$



Notation

■ Numbers:

- N : total number of samples
- M : model order, number of features (engineered or not)
- K : number of outputs or classes

■ Vectors:

- \mathbf{x} : feature vector, $\mathbf{x} = [x_1, x_2, \dots, x_M]^T$
- \mathbf{y} : target vector, $\mathbf{y} = [y_1, y_2, \dots, y_K]^T$
- $\hat{\mathbf{y}}$: predicted target vector, $\hat{\mathbf{y}} = [\hat{y}_1, \hat{y}_2, \dots, \hat{y}_K]^T$
- \mathbf{w} : weight vector for $K = 1$ targets, $\mathbf{w} = [w_1, w_2, \dots, w_M]^T$
 - Also used: $\boldsymbol{\beta} = [\beta_1, \beta_2, \dots, \beta_M]^T$
- \mathbf{b} : bias vector, $\mathbf{b} = [b_1, b_2, \dots, b_K]$

Notation

■ Numbers:

- N : total number of samples
- M : model order, number of features (engineered or not)
- K : number of outputs or classes

■ Vectors:

- \mathbf{x} : feature vector, $\mathbf{x} = [x_1, x_2, \dots, x_M]^T$
- \mathbf{y} : target vector, $\mathbf{y} = [y_1, y_2, \dots, y_K]^T$
- $\hat{\mathbf{y}}$: predicted target vector, $\hat{\mathbf{y}} = [\hat{y}_1, \hat{y}_2, \dots, \hat{y}_K]^T$
- \mathbf{w} : weight vector for $K = 1$ targets, $\mathbf{w} = [w_1, w_2, \dots, w_M]^T$
 - Also used: $\boldsymbol{\beta} = [\beta_1, \beta_2, \dots, \beta_M]^T$
- \mathbf{b} : bias vector, $\mathbf{b} = [b_1, b_2, \dots, b_K]$
 - Also used: β_0 for $K = 1$

Notation

- Numbers:
 - N : total number of samples
 - M : model order, number of features (engineered or not)
 - K : number of outputs or classes
- Vectors:
 - \mathbf{x} : feature vector, $\mathbf{x} = [x_1, x_2, \dots, x_M]^T$
 - \mathbf{y} : target vector, $\mathbf{y} = [y_1, y_2, \dots, y_K]^T$
 - $\hat{\mathbf{y}}$: predicted target vector, $\hat{\mathbf{y}} = [\hat{y}_1, \hat{y}_2, \dots, \hat{y}_K]^T$
 - \mathbf{w} : weight vector for $K = 1$ targets, $\mathbf{w} = [w_1, w_2, \dots, w_M]^T$
 - Also used: $\beta = [\beta_1, \beta_2, \dots, \beta_M]^T$
 - \mathbf{b} : bias vector, $\mathbf{b} = [b_1, b_2, \dots, b_K]$
 - Also used: β_0 for $K = 1$
- Matrices:

Notation

■ Numbers:

- N : total number of samples
- M : model order, number of features (engineered or not)
- K : number of outputs or classes

■ Vectors:

- \mathbf{x} : feature vector, $\mathbf{x} = [x_1, x_2, \dots, x_M]^T$
- \mathbf{y} : target vector, $\mathbf{y} = [y_1, y_2, \dots, y_K]^T$
- $\hat{\mathbf{y}}$: predicted target vector, $\hat{\mathbf{y}} = [\hat{y}_1, \hat{y}_2, \dots, \hat{y}_K]^T$
- \mathbf{w} : weight vector for $K = 1$ targets, $\mathbf{w} = [w_1, w_2, \dots, w_M]^T$
 - Also used: $\boldsymbol{\beta} = [\beta_1, \beta_2, \dots, \beta_M]^T$
- \mathbf{b} : bias vector, $\mathbf{b} = [b_1, b_2, \dots, b_K]$
 - Also used: β_0 for $K = 1$

■ Matrices:

- A : (N, M) design matrix

Notation

■ Numbers:

- N : total number of samples
- M : model order, number of features (engineered or not)
- K : number of outputs or classes

■ Vectors:

- \mathbf{x} : feature vector, $\mathbf{x} = [x_1, x_2, \dots, x_M]^T$
- \mathbf{y} : target vector, $\mathbf{y} = [y_1, y_2, \dots, y_K]^T$
- $\hat{\mathbf{y}}$: predicted target vector, $\hat{\mathbf{y}} = [\hat{y}_1, \hat{y}_2, \dots, \hat{y}_K]^T$
- \mathbf{w} : weight vector for $K = 1$ targets, $\mathbf{w} = [w_1, w_2, \dots, w_M]^T$
 - Also used: $\beta = [\beta_1, \beta_2, \dots, \beta_M]^T$
- \mathbf{b} : bias vector, $\mathbf{b} = [b_1, b_2, \dots, b_K]$
 - Also used: β_0 for $K = 1$

■ Matrices:

- A : (N, M) design matrix
 - Also used: X

Notation

■ Numbers:

- N : total number of samples
- M : model order, number of features (engineered or not)
- K : number of outputs or classes

■ Vectors:

- \mathbf{x} : feature vector, $\mathbf{x} = [x_1, x_2, \dots, x_M]^T$
- \mathbf{y} : target vector, $\mathbf{y} = [y_1, y_2, \dots, y_K]^T$
- $\hat{\mathbf{y}}$: predicted target vector, $\hat{\mathbf{y}} = [\hat{y}_1, \hat{y}_2, \dots, \hat{y}_K]^T$
- \mathbf{w} : weight vector for $K = 1$ targets, $\mathbf{w} = [w_1, w_2, \dots, w_M]^T$
 - Also used: $\boldsymbol{\beta} = [\beta_1, \beta_2, \dots, \beta_M]^T$
- \mathbf{b} : bias vector, $\mathbf{b} = [b_1, b_2, \dots, b_K]$
 - Also used: β_0 for $K = 1$

■ Matrices:

- A : (N, M) design matrix
 - Also used: X
- W : (K, M) weight matrix

Supervised Learning

Type	Linear Regression	Logistic Regression
Use	Modeling Continuous Data	Classification
Features	Any Numerical Data, $\mathbf{x} = [x_1, x_2, \dots, x_M]^T$	
Targets	Any Numerical Data, \mathbf{y}	Class Labels, \mathbf{y}
Model	$\hat{\mathbf{y}} = \mathbf{W}\mathbf{x} + \mathbf{b}$	$\hat{\mathbf{y}} = \sigma(\mathbf{W}\mathbf{x} + \mathbf{b})$
Loss Function	Error between \mathbf{y} and $\hat{\mathbf{y}}$	Cross-Entropy

Optimization

- Use loss/error/cost function to find best model-parameters

Problem	Loss Function	Formula
Regression	Squared/L2 Loss	$\sum_i (\mathbf{y}_i - \hat{\mathbf{y}}_i)^2$
Binary Classification	Binary Cross-Entropy	$-\sum_i (y_i \ln(\hat{y}_i) + (1 - y_i) \ln(1 - \hat{y}_i))$
Multi-Class Classification	Cross-Entropy	$-\sum_i \sum_k (y_{ik} \ln(\hat{y}_{ik}))$

Optimization

- Use loss/error/cost function to find best model-parameters
- Non-linear opt. can use arbitrary Loss function

Problem	Loss Function	Formula
Regression	Squared/L2 Loss	$\sum_i (\mathbf{y}_i - \hat{\mathbf{y}}_i)^2$
Binary Classification	Binary Cross-Entropy	$-\sum_i (y_i \ln(\hat{y}_i) + (1 - y_i) \ln(1 - \hat{y}_i))$
Multi-Class Classification	Cross-Entropy	$-\sum_i \sum_k (y_{ik} \ln(\hat{y}_{ik}))$

Goodness of Fit

- Evaluate the accuracy of the model

Goodness of Fit

- Evaluate the accuracy of the model
- Can use criteria different than that used for optimization

Goodness of Fit

- Evaluate the accuracy of the model
- Can use criteria different than that used for optimization
- Examples:

Goodness of Fit

- Evaluate the accuracy of the model
- Can use criteria different than that used for optimization
- Examples:
 - Mean Squared Error: $\frac{1}{N} \sum_i (\mathbf{y}_i - \hat{\mathbf{y}}_i)^2$

Goodness of Fit

- Evaluate the accuracy of the model
- Can use criteria different than that used for optimization
- Examples:
 - Mean Squared Error: $\frac{1}{N} \sum_i (\mathbf{y}_i - \hat{\mathbf{y}}_i)^2$
 - May also represent result of optimization

Goodness of Fit

- Evaluate the accuracy of the model
- Can use criteria different than that used for optimization
- Examples:
 - Mean Squared Error: $\frac{1}{N} \sum_i (\mathbf{y}_i - \hat{\mathbf{y}}_i)^2$
 - May also represent result of optimization
 - Mean Absolute Error: $\frac{1}{N} \sum_i |\mathbf{y}_i - \hat{\mathbf{y}}_i|$

Goodness of Fit

- Evaluate the accuracy of the model
- Can use criteria different than that used for optimization
- Examples:
 - Mean Squared Error: $\frac{1}{N} \sum_i (\mathbf{y}_i - \hat{\mathbf{y}}_i)^2$
 - May also represent result of optimization
 - Mean Absolute Error: $\frac{1}{N} \sum_i |\mathbf{y}_i - \hat{\mathbf{y}}_i|$
 - Easily interpretable units

Goodness of Fit

- Evaluate the accuracy of the model
- Can use criteria different than that used for optimization
- Examples:
 - Mean Squared Error: $\frac{1}{N} \sum_i (\mathbf{y}_i - \hat{\mathbf{y}}_i)^2$
 - May also represent result of optimization
 - Mean Absolute Error: $\frac{1}{N} \sum_i |\mathbf{y}_i - \hat{\mathbf{y}}_i|$
 - Easily interpretable units
 - Root Mean Squared Error: $\sqrt{\frac{1}{N} \sum_i (\mathbf{y}_i - \hat{\mathbf{y}}_i)^2}$

Goodness of Fit

- Evaluate the accuracy of the model
- Can use criteria different than that used for optimization
- Examples:
 - Mean Squared Error: $\frac{1}{N} \sum_i (\mathbf{y}_i - \hat{\mathbf{y}}_i)^2$
 - May also represent result of optimization
 - Mean Absolute Error: $\frac{1}{N} \sum_i |\mathbf{y}_i - \hat{\mathbf{y}}_i|$
 - Easily interpretable units
 - Root Mean Squared Error: $\sqrt{\frac{1}{N} \sum_i (\mathbf{y}_i - \hat{\mathbf{y}}_i)^2}$
 - May represent opt. & easily interpretable units

Goodness of Fit


- Evaluate the accuracy of the model
- Can use criteria different than that used for optimization
- Examples:
 - Mean Squared Error: $\frac{1}{N} \sum_i (\mathbf{y}_i - \hat{\mathbf{y}}_i)^2$
 - May also represent result of optimization
 - Mean Absolute Error: $\frac{1}{N} \sum_i |\mathbf{y}_i - \hat{\mathbf{y}}_i|$
 - Easily interpretable units
 - Root Mean Squared Error: $\sqrt{\frac{1}{N} \sum_i (\mathbf{y}_i - \hat{\mathbf{y}}_i)^2}$
 - May represent opt. & easily interpretable units
 - Classification Accuracy: $\frac{1}{N} \sum_i (\mathbf{y}_i == \hat{\mathbf{y}}_i)$

Train, Validation, and Test Sets

- Always split your data into train and test sets to see how well it does against new data
- Train set: set of data to be used for training
 - e.g. `model.fit(x_train,y_train)`
- Test set: After training is done, evaluate how well it does against unseen data using test set
- Validation set: If tuning hyper-paramters, perform one more split to get a validation set. Use validation set to tune parameters

Train and Test Sets (Dealing with Time Series)

- Train and test split is usually done by taking samples at random from the entire data set
- But when using time series to predict future, it is better to select test set to be a continuous chunk at the end of the time series
- Because we want to see how well the model does in predicting the future



Regularization

- Prevent over-fitting by adding a term to loss function
- *Loss Function = Target loss function + λ Regularization*
- λ hyper-parameter determine how much to emphasize on regularizing
- Large weights usually lead to over-fitting
- Weight-based regularization is most commonly used
 - L2 (Ridge) Regularization: $\sum_{j=1}^D |w_j|^2$
 - L1 (Lasso) Regularization: $\sum_{j=1}^D |w_j|$
- First over-estimate the model order you need, then use regularization to prevent over-fitting

Outline

- 1 Review of Week 1
- 2 Neural Network Model
- 3 Training with Neural Networks
- 4 Introduction to Keras
- 5 Lab: Music Classification
- 6 (Optional) Lab: Cat vs. Non-Cat

Extension from Logistic Regression

- Logistic Regression Model: $\hat{y} = \sigma(W\mathbf{x} + b)$

Extension from Logistic Regression

- Logistic Regression Model: $\hat{y} = \sigma(W\mathbf{x} + b)$
- Replace \mathbf{x} with $\mathbf{z} = f(\mathbf{x})$: $\hat{y} = \sigma(W\mathbf{z} + b)$

Extension from Logistic Regression

- Logistic Regression Model: $\hat{y} = \sigma(W\mathbf{x} + b)$
- Replace \mathbf{x} with $\mathbf{z} = f(\mathbf{x})$: $\hat{y} = \sigma(W\mathbf{z} + b)$
 - Think of feature engineering...

Extension from Logistic Regression

- Logistic Regression Model: $\hat{y} = \sigma(W\mathbf{x} + b)$
- Replace \mathbf{x} with $\mathbf{z} = f(\mathbf{x})$: $\hat{y} = \sigma(W\mathbf{z} + b)$
 - Think of feature engineering...
- What if $f(\mathbf{x})$ is a linear transformation? ie. $f(\mathbf{x}) = W'\mathbf{x}$

Extension from Logistic Regression

- Logistic Regression Model: $\hat{y} = \sigma(W\mathbf{x} + b)$
- Replace \mathbf{x} with $\mathbf{z} = f(\mathbf{x})$: $\hat{y} = \sigma(W\mathbf{z} + b)$
 - Think of feature engineering...
- What if $f(\mathbf{x})$ is a linear transformation? ie. $f(\mathbf{x}) = W'\mathbf{x}$
 - All linear transforms can be represented as matrices

Extension from Logistic Regression

- Logistic Regression Model: $\hat{y} = \sigma(W\mathbf{x} + b)$
- Replace \mathbf{x} with $\mathbf{z} = f(\mathbf{x})$: $\hat{y} = \sigma(W\mathbf{z} + b)$
 - Think of feature engineering...
- What if $f(\mathbf{x})$ is a linear transformation? ie. $f(\mathbf{x}) = W'\mathbf{x}$
 - All linear transforms can be represented as matrices
 - So, $\hat{y} = \sigma(W''\mathbf{x} + b)$, where $W'' = WW'$

Extension from Logistic Regression

- Logistic Regression Model: $\hat{y} = \sigma(W\mathbf{x} + b)$
- Replace \mathbf{x} with $\mathbf{z} = f(\mathbf{x})$: $\hat{y} = \sigma(W\mathbf{z} + b)$
 - Think of feature engineering...
- What if $f(\mathbf{x})$ is a linear transformation? ie. $f(\mathbf{x}) = W'\mathbf{x}$
 - All linear transforms can be represented as matrices
 - So, $\hat{y} = \sigma(W''\mathbf{x} + b)$, where $W'' = WW'$
- Non-linear transform gives something different...

Extension from Logistic Regression

- Logistic Regression Model: $\hat{y} = \sigma(W\mathbf{x} + b)$
- Replace \mathbf{x} with $\mathbf{z} = f(\mathbf{x})$: $\hat{y} = \sigma(W\mathbf{z} + b)$
 - Think of feature engineering...
- What if $f(\mathbf{x})$ is a linear transformation? ie. $f(\mathbf{x}) = W'\mathbf{x}$
 - All linear transforms can be represented as matrices
 - So, $\hat{y} = \sigma(W''\mathbf{x} + b)$, where $W'' = WW'$
- Non-linear transform gives something different...
 - Recall polynomial transformations and exponential transformations of the data

Extension from Logistic Regression

- Logistic Regression Model: $\hat{y} = \sigma(W\mathbf{x} + b)$
- Replace \mathbf{x} with $\mathbf{z} = f(\mathbf{x})$: $\hat{y} = \sigma(W\mathbf{z} + b)$
 - Think of feature engineering...
- What if $f(\mathbf{x})$ is a linear transformation? ie. $f(\mathbf{x}) = W'\mathbf{x}$
 - All linear transforms can be represented as matrices
 - So, $\hat{y} = \sigma(W''\mathbf{x} + b)$, where $W'' = WW'$
- Non-linear transform gives something different...
 - Recall polynomial transformations and exponential transformations of the data
 - These cannot be expressed as matrix multiplication

Extension to Neural Network


- Restrict $f(\mathbf{x})$ to non-linear function applied to all input values

Extension to Neural Network

- Restrict $f(\mathbf{x})$ to non-linear function applied to all input values
 - Simplest example of a **Neural Network**


Extension to Neural Network

- Restrict $f(\mathbf{x})$ to non-linear function applied to all input values
 - Simplest example of a **Neural Network**
- $\hat{y} = \sigma(W_2 f_1(W_1 \mathbf{x} + \mathbf{b}_1) + b_2)$




Extension to Neural Network

- Restrict $f(\mathbf{x})$ to non-linear function applied to all input values
 - Simplest example of a **Neural Network**
- $\hat{y} = \sigma(W_2 f_1(W_1 \mathbf{x} + \mathbf{b}_1) + b_2)$
- We can optimize for both W_1, \mathbf{b}_1 and W_2, b_2 model-parameters




Extension to Neural Network

- Restrict $f(\mathbf{x})$ to non-linear function applied to all input values
 - Simplest example of a **Neural Network**
- $\hat{y} = \sigma(W_2 f_1(W_1 \mathbf{x} + \mathbf{b}_1) + b_2)$
- We can optimize for both W_1, \mathbf{b}_1 and W_2, b_2 model-parameters
 - $\nabla J = [\frac{\partial J}{\partial w_1}, \frac{\partial J}{\partial w_2}, \dots, \frac{\partial J}{\partial w_p}]^T$




Extension to Neural Network

- Restrict $f(\mathbf{x})$ to non-linear function applied to all input values
 - Simplest example of a **Neural Network**
- $\hat{y} = \sigma(W_2 f_1(W_1 \mathbf{x} + \mathbf{b}_1) + b_2)$
- We can optimize for both W_1, \mathbf{b}_1 and W_2, b_2 model-parameters
 - $\nabla J = [\frac{\partial J}{\partial w_1}, \frac{\partial J}{\partial w_2}, \dots, \frac{\partial J}{\partial w_p}]^T$
 - Now we're *learning* the feature engineering



Extension to Neural Network

- Restrict $f(\mathbf{x})$ to non-linear function applied to all input values
 - Simplest example of a **Neural Network**
- $\hat{y} = \sigma(W_2 f_1(W_1 \mathbf{x} + \mathbf{b}_1) + b_2)$
- We can optimize for both W_1, \mathbf{b}_1 and W_2, b_2 model-parameters
 - $\nabla J = [\frac{\partial J}{\partial w_1}, \frac{\partial J}{\partial w_2}, \dots, \frac{\partial J}{\partial w_p}]^T$
 - Now we're *learning* the feature engineering
- But why stop here?...



Mathematical Model: Multi-Layer Perceptron

■ Model:

$$\hat{\mathbf{y}} = f_{out}(f_L(W_L \mathbf{z}_L + b_L)) + b_{out}$$

Mathematical Model: Multi-Layer Perceptron

- Model:

$$\hat{\mathbf{y}} = f_{out}(f_L(W_L \mathbf{z}_L + b_L)) + b_{out}$$

- Where, $\mathbf{z}_l = f_{l-1}(W_{l-1} \mathbf{z}_{l-1} + b_{l-1})$ for $1 \leq l \leq L$, $\mathbf{z}_0 := \mathbf{x}$, and L is the number of hidden layers

Mathematical Model: Multi-Layer Perceptron

- Model:

$$\hat{\mathbf{y}} = f_{out}(f_L(W_L \mathbf{z}_L + b_L)) + b_{out})$$

- Where, $\mathbf{z}_l = f_{l-1}(W_{l-1} \mathbf{z}_{l-1} + b_{l-1})$ for $1 \leq l \leq L$, $\mathbf{z}_0 := \mathbf{x}$, and L is the number of hidden layers
- ie. all hidden layers are non-linear activation of linear transform

Mathematical Model: Multi-Layer Perceptron

■ Model:

$$\hat{\mathbf{y}} = f_{out}(f_L(W_L \mathbf{z}_L + b_L)) + b_{out}$$

- Where, $\mathbf{z}_l = f_{l-1}(W_{l-1} \mathbf{z}_{l-1} + b_{l-1})$ for $1 \leq l \leq L$, $\mathbf{z}_0 := \mathbf{x}$, and L is the number of hidden layers
- ie. all hidden layers are non-linear activation of linear transform
- f_{out} depends on type of ML problem (regression: linear, classification: sigmoid/soft-max)

Mathematical Model: Multi-Layer Perceptron

■ Model:

$$\hat{\mathbf{y}} = f_{out}(f_L(W_L \mathbf{z}_L + b_L)) + b_{out}$$

- Where, $\mathbf{z}_l = f_{l-1}(W_{l-1} \mathbf{z}_{l-1} + b_{l-1})$ for $1 \leq l \leq L$, $\mathbf{z}_0 := \mathbf{x}$, and L is the number of hidden layers
- ie. all hidden layers are non-linear activation of linear transform
- f_{out} depends on type of ML problem (regression: linear, classification: sigmoid/soft-max)
- Layers: Input, Hidden, Output

Mathematical Model: Multi-Layer Perceptron

■ Model:

$$\hat{\mathbf{y}} = f_{out}(f_L(W_L \mathbf{z}_L + b_L)) + b_{out}$$

- Where, $\mathbf{z}_l = f_{l-1}(W_{l-1}\mathbf{z}_{l-1} + b_{l-1})$ for $1 \leq l \leq L$, $\mathbf{z}_0 := \mathbf{x}$, and L is the number of hidden layers
- ie. all hidden layers are non-linear activation of linear transform
- f_{out} depends on type of ML problem (regression: linear, classification: sigmoid/soft-max)
- Layers: Input, Hidden, Output
 - All layers before output performing feature extraction for final layer linear/logistic regression

Mathematical Model: Multi-Layer Perceptron

■ Model:

$$\hat{\mathbf{y}} = f_{out}(f_L(W_L \mathbf{z}_L + b_L)) + b_{out}$$


- Where, $\mathbf{z}_l = f_{l-1}(W_{l-1}\mathbf{z}_{l-1} + b_{l-1})$ for $1 \leq l \leq L$, $\mathbf{z}_0 := \mathbf{x}$, and L is the number of hidden layers
- ie. all hidden layers are non-linear activation of linear transform
- f_{out} depends on type of ML problem (regression: linear, classification: sigmoid/soft-max)
- Layers: Input, Hidden, Output
 - All layers before output performing feature extraction for final layer linear/logistic regression
 - layer size determined by matrix multiply W_l (output dim, input dim)

Mathematical Model: Multi-Layer Perceptron

- Activation Functions: On board

Toy Example: Spiral Classification


Human Engineered
Feature Transformations:




NN Engineered
Feature Transformations:

Toy Example: Spiral Classification

Human Engineered
Feature Transformations:



NN Engineered
Feature Transformations:



Advantages and Disadvantages

Advantages

- Further removed need for domain knowledge
- Infinitely expressive

Disadvantages

- Less control over behavior of model
- Computationally expensive

Biological Justification


- Example: Steps for Processing Vision
 - 1 Eyes gather light
 - 2 Light intensities converted to shapes
 - 3 shapes recognized as objects

Outline

- 1 Review of Week 1
- 2 Neural Network Model
- 3 Training with Neural Networks
- 4 Introduction to Keras
- 5 Lab: Music Classification
- 6 (Optional) Lab: Cat vs. Non-Cat

Large Scale Machine Learning



- Learning with large data sets
- Algorithms today perform so much better than five years ago due to sheer amount of data availability
- “It’s not who has the best algorithm that wins. It’s who has the most data”
 - So we want to learn from large data sets



Learning with Large Data Sets

- Challenges:
 - Computationally very expensive to compute gradients
 - And each gradient computation performs only one step of update
- In large scale machine learning, we want to come up with computationally reasonable ways to deal with large data sets
 - Batch Gradient Descent
 - Stochastic Gradient Descent
 - Mini-batch Gradient Descent

Digression: Revisiting Learning Rate



Batch Gradient Descent

- **Batch Gradient Descent** takes all the examples in the training data to compute one step of gradient descent update
- Algorithm: Consider linear regression ($N = 100,000,000$)
 - $\hat{y} = \sum_{i=0}^N w_i x_i$
 - Cost, $J = \frac{1}{N} \sum_{i=0}^N (y_i - \hat{y}_i)^2$
 - Gradient Descent Update $w_{new} = w_{old} - \alpha \frac{dJ}{dw}$

Stochastic Gradient Descent

- **SGD** takes only one example in the training example to perform one step of gradient descent
 - The algorithm modifies the parameters a little bit to fit the just first example (x_1, y_1)
 - Then again modify the parameters to fit the second training example (x_2, y_2) and so on...
- Algorithm (Let N be the total number of training examples):
Repeat{
 for $i = 1, 2 \dots N$ {
 $Cost, J = (y_i - \hat{y}_i)^2$
 Gradient Descent Update $w_{new} = w_{old} - \alpha \frac{dJ}{dw}$
 }
}

Batch Gradient Descent

- **Batch Gradient Descent** uses 'b' training examples to perform one update step

- 'b' is called batch size
- Number of iterations = $\frac{N}{b}$

- Algorithm:

Repeat{

$j = 0$

for i in range(iterations){

$$Cost, J = \frac{1}{b} \sum_{j=1}^{i+b} (y_j - \hat{y}_j)^2$$

Gradient Descent Update $w_{new} = w_{old} - \alpha \frac{dJ}{dw}$

$j = j + b$

}

}

Outline

- 1 Review of Week 1
- 2 Neural Network Model
- 3 Training with Neural Networks
- 4 Introduction to Keras
- 5 Lab: Music Classification
- 6 (Optional) Lab: Cat vs. Non-Cat

Outline

- 1 Review of Week 1
- 2 Neural Network Model
- 3 Training with Neural Networks
- 4 Introduction to Keras
- 5 Lab: Music Classification**
- 6 (Optional) Lab: Cat vs. Non-Cat

Outline

- 1 Review of Week 1
- 2 Neural Network Model
- 3 Training with Neural Networks
- 4 Introduction to Keras
- 5 Lab: Music Classification
- 6 (Optional) Lab: Cat vs. Non-Cat

Learning Objectives

- What are the advantages of Neural Networks?
- What is the mathematical model for a Neural Network?
- What are the hyper-parameters associated with a Neural Network?
- Why are batch-size and learning-rate important? How are they related?
- How do we implement Neural Networks with Keras?

Thank You!

- Next Class: Convolutional Neural Networks