Day 10: Final Project Presentations & Going Further with ML/DL

Summer STEM: Machine Learning

Nikola Janjušević, Akshaj Kumar Veldanda, Jacky Yuan, Tejaishwarya Gagadam

Department of Electrical and Computer Engineering
NYU Tandon School of Engineering
Brooklyn, New York

June 28, 2019

Outline

1 Schedule

2 Presentation Guidelines

3 Akshaj's Talk

4 Generative Adversarial Networks

5 Course Takeaway

Schedule

1. Finishing Up Projects (9:00-11:00am)
2. Make Presentations (11:00-12:00am)
3. Lunch (12:00pm-1:00pm)
4. Presentations (1:10-3:00pm)
5. Final Material: Going Further with ML/DL (3:00-4:00pm)
6. Goodbye! (4:00pm - ∞?)

Outline

1. Schedule
2. Presentation Guidelines
3. Akshaj's Talk
4. Generative Adversarial Networks
5. Course Takeaway

Presentation Guidelines

- 8-10 minutes per group
- Talk about your project:
  - **Goal**: features and targets. What are you predicting?
  - **Dataset**: Where is it from, what does it consist of. Features, targets.
  - **Model**: What is your model? MLP? CNN? What is the architecture? Why?
  - **Training**: How did your model train? Challenges?
  - **Inference**: How was your models predictions on new data?
  - **Going Forward**: "We'd like to continue by..."
  - **Learning**: what did you learn/take away from this project? From this course?
- Put your slides on the master Google-Slides slide-deck (check email)

NYU TANDON SCHOOL OF ENGINEERING

Outline

1 Schedule

2 Presentation Guidelines

3 Akshaj's Talk

4 Generative Adversarial Networks

5 Course Takeaway

Outline

1 Schedule

2 Presentation Guidelines

3 Akshaj's Talk

4 Generative Adversarial Networks

5 Course Takeaway

GANs: Generative Adversarial Networks

- Invented in 2014 by Ian Goodfellow

GANs: Generative Adversarial Networks

- Invented in 2014 by Ian Goodfellow
- Goal: generate samples never seen before

GANs: Generative Adversarial Networks

- Invented in 2014 by Ian Goodfellow
- Goal: generate samples never seen before
- How: game between two networks

GANs: Generative Adversarial Networks

- Invented in 2014 by Ian Goodfellow
- Goal: generate samples never seen before
- How: game between two networks
  - Generator Network

GANs: Generative Adversarial Networks

- Invented in 2014 by Ian Goodfellow
- Goal: generate samples never seen before
- How: game between two networks
  - Generator Network
  - Discriminator Network

GANs: Generative Adversarial Networks

- Invented in 2014 by Ian Goodfellow
- Goal: generate samples never seen before
- How: game between two networks
  - Generator Network
  - Discriminator Network
- **Goal of Generator**: generate fake samples indistinguishable from real samples

GANs: Generative Adversarial Networks

- Invented in 2014 by Ian Goodfellow
- Goal: generate samples never seen before
- How: game between two networks
  - Generator Network
  - Discriminator Network
- **Goal of Generator**: generate fake samples indistinguishable from real samples
- **Goal of Discriminator**: be able to tell apart real and fake samples

Implementation

- Generator Model: (on board)

Implementation

- Generator Model: (on board)
  - Draw sample from multi-dim noise vector

Implementation

- Generator Model: (on board)
  - Draw sample from multi-dim noise vector
  - Up-Convolve into images (learning kernels to generate features)

Implementation

- Generator Model: (on board)
  - Draw sample from multi-dim noise vector
  - Up-Convolve into images (learning kernels to generate features)
- Discriminator Model: (on board)

Implementation

- Generator Model: (on board)
  - Draw sample from multi-dim noise vector
  - Up-Convolve into images (learning kernels to generate features)
- Discriminator Model: (on board)
  - Binary Classifier (CNN or MLP)

Implementation

- Generator Model: (on board)
  - Draw sample from multi-dim noise vector
  - Up-Convolve into images (learning kernels to generate features)
- Discriminator Model: (on board)
  - Binary Classifier (CNN or MLP)
- Training:

Implementation

- Generator Model: (on board)
  - Draw sample from multi-dim noise vector
  - Up-Convolve into images (learning kernels to generate features)
- Discriminator Model: (on board)
  - Binary Classifier (CNN or MLP)
- Training:
  - Randomly feed Discriminator real and generated samples

Implementation

- Generator Model: (on board)
  - Draw sample from multi-dim noise vector
  - Up-Convolve into images (learning kernels to generate features)
- Discriminator Model: (on board)
  - Binary Classifier (CNN or MLP)
- Training:
  - Randomly feed Discriminator real and generated samples
  - Balancing act: train D, train G, train D, train G, ...

Implementation

- Generator Model: (on board)
  - Draw sample from multi-dim noise vector
  - Up-Convolve into images (learning kernels to generate features)
- Discriminator Model: (on board)
  - Binary Classifier (CNN or MLP)
- Training:
  - Randomly feed Discriminator real and generated samples
  - Balancing act: train D, train G, train D, train G, ...
  - Cost function for each network?

Implementation

- Generator Model: (on board)
  - Draw sample from multi-dim noise vector
  - Up-Convolve into images (learning kernels to generate features)
- Discriminator Model: (on board)
  - Binary Classifier (CNN or MLP)
- Training:
  - Randomly feed Discriminator real and generated samples
  - Balancing act: train D, train G, train D, train G, ...
  - Cost function for each network?
    - Binary Cross Entropy for each

Implementation

- Generator Model: (on board)
  - Draw sample from multi-dim noise vector
  - Up-Convolve into images (learning kernels to generate features)
- Discriminator Model: (on board)
  - Binary Classifier (CNN or MLP)
- Training:
  - Randomly feed Discriminator real and generated samples
  - Balancing act: train D, train G, train D, train G, ...
  - Cost function for each network?
    - Binary Cross Entropy for each
    - Freeze D: Update weights of G to **maximize** BCE

Implementation

- Generator Model: (on board)
  - Draw sample from multi-dim noise vector
  - Up-Convolve into images (learning kernels to generate features)
- Discriminator Model: (on board)
  - Binary Classifier (CNN or MLP)
- Training:
  - Randomly feed Discriminator real and generated samples
  - Balancing act: train D, train G, train D, train G, ...
  - Cost function for each network?
    - Binary Cross Entropy for each
    - Freeze D: Update weights of G to **maximize** BCE
    - Freeze G: Update weights of D to **minimize** BCE

Implementation

- Generator Model: (on board)
  - Draw sample from multi-dim noise vector
  - Up-Convolve into images (learning kernels to generate features)
- Discriminator Model: (on board)
  - Binary Classifier (CNN or MLP)
- Training:
  - Randomly feed Discriminator real and generated samples
  - Balancing act: train D, train G, train D, train G, ...
  - Cost function for each network?
    - Binary Cross Entropy for each
    - Freeze D: Update weights of G to **maximize** BCE
    - Freeze G: Update weights of D to **minimize** BCE
  - Game-theory: converge to maximin of the cost function

Implementation

- Generator Model: (on board)
  - Draw sample from multi-dim noise vector
  - Up-Convolve into images (learning kernels to generate features)
- Discriminator Model: (on board)
  - Binary Classifier (CNN or MLP)
- Training:
  - Randomly feed Discriminator real and generated samples
  - Balancing act: train D, train G, train D, train G, ...
  - Cost function for each network?
    - Binary Cross Entropy for each
    - Freeze D: Update weights of G to **maximize** BCE
    - Freeze G: Update weights of D to **minimize** BCE
  - Game-theory: converge to maximin of the cost function
- Multi-dim noise vector allows interpolation between samples and can be trained to determine output class of image

Beginning

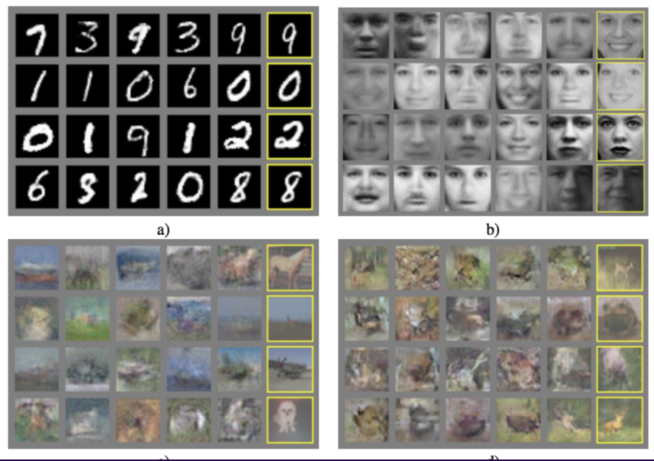

Generated images (yellow) on a) MNIST b) TFD c) CIFAR-10 (MLP model) d) CIFAR-10 (Conv model)
"Generative Adversarial Networks", Goodfellow et. al. 2014

Progress

Improvement of GANs in producing photo-realistic faces over the years

Celebrity Faces

Human face generation one of the most difficult tasks

Occlusion

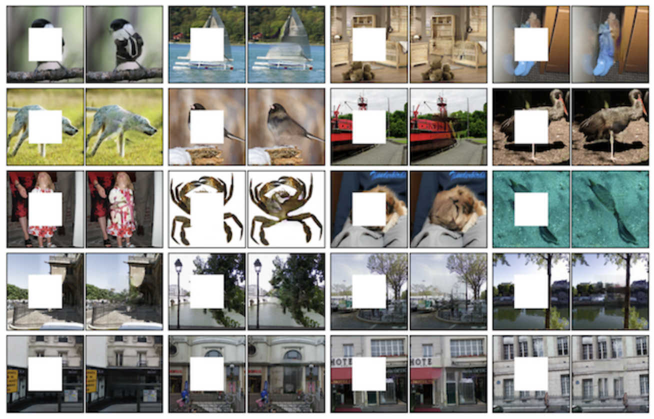

Recreate missing parts of an image

Text to Image

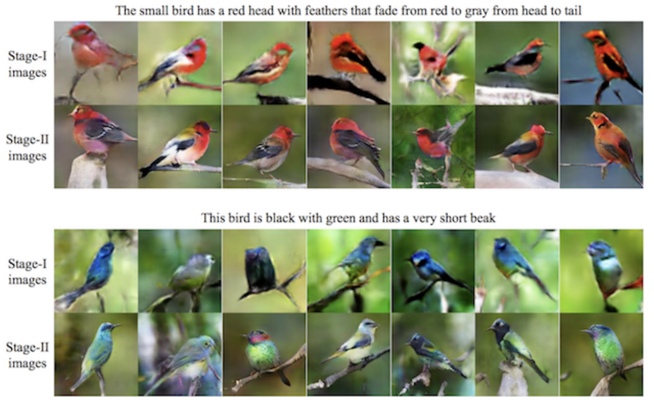

produce image based on a description

Videos

- Progressive Growing of GANs, 2017
- Few-Shot Adversarial Learning of Realistic Neural Talking Head Models, 2019

Applications of GANs?

- Applications often obscure/tough to think of
- Entering realm of computers outperforming humans on creative tasks

Outline

1. Schedule
2. Presentation Guidelines
3. Akshaj's Talk
4. Generative Adversarial Networks
5. Course Takeaway

Course Takeaway

- ML is a lot of Math Math Math Math

Course Takeaway

- ML is a lot of Math Math Math Math
- We've only shown you a subsection

Course Takeaway

- ML is a lot of Math Math Math Math
- We've only shown you a subsection
  - Supervised Learning: Linear/Logistic Regression and Neural Networks

Course Takeaway

- ML is a lot of Math Math Math Math
- We've only shown you a subsection
  - Supervised Learning: Linear/Logistic Regression and Neural Networks
- If you're interested in pursuing further, you know where to look!

Course Takeaway

- ML is a lot of Math Math Math Math
- We've only shown you a subsection
  - Supervised Learning: Linear/Logistic Regression and Neural Networks
- If you're interested in pursuing further, you know where to look!
- If not, that's valuable information for deciding what to do with your future!

Thank You!

- Thank You, Thank You, Thank You!